

Day 6: Principles of Network Applications



CSEE 4119
Computer Networks
Ethan Katz-Bassett

 COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

Slides adapted from (and often identical to) slides from Kurose and Ross.

All material copyright 1996-2020

J.F.Kurose and K.W.Ross, All Rights Reserved

Sept 22 admin

- Masks **required**, over nose and mouth
- If you are not feeling well or were exposed to COVID, please stay home and watch video
 - No more Zoom option
- HWI, preliminary stage of Project I to be released soon
- Review socket programming on your own
 - 1:10 video from today by Shuyue (TA)
 - “Shuyue was spectacular”
 - material in the book
 - online resources
 - just try it! (with the first part of Project I)
 - attend office hours with questions

Chapter 2: outline

Today's class

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

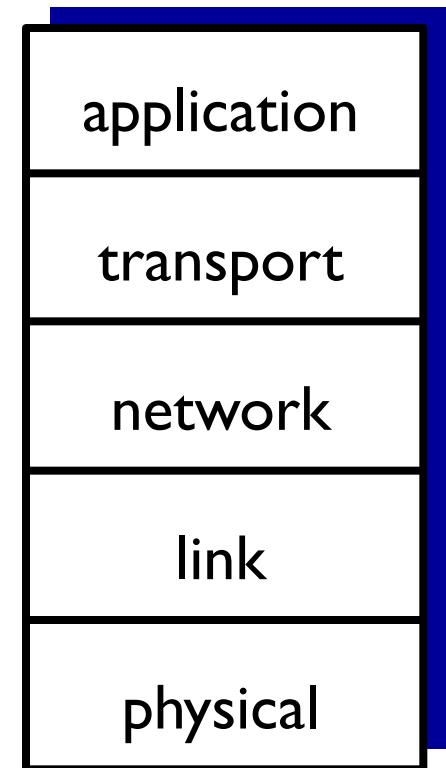
2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Recap: Internet protocol stack

- ***application***: supporting network applications
 - FTP, SMTP, HTTP *This chapter*
- ***transport***: process-process data transfer
 - TCP, UDP
- ***network***: routing of datagrams from source to destination
 - IP, routing protocols
- ***link***: data transfer between neighboring network elements
 - Ethernet, 802.111 (WiFi), PPP
- ***physical***: bits “on the wire”



Chapter 2: application layer

our goals:

- conceptual, implementation aspects of network application protocols
 - transport-layer service models
 - client-server paradigm
 - peer-to-peer paradigm
 - content distribution networks
- learn about protocols by examining popular application-level protocols
 - HTTP
 - SMTP / POP3 / IMAP
 - DNS
 - video streaming
- creating network applications
 - socket API

Some network apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

Creating a network app

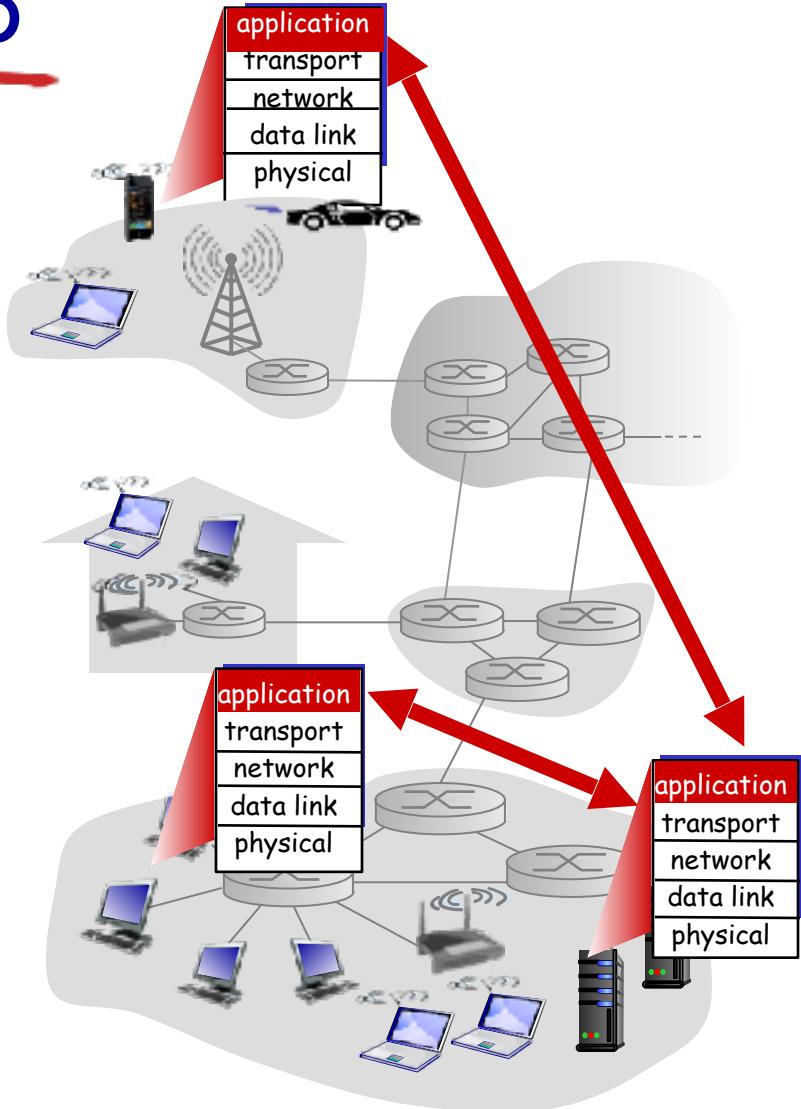
You'll do this in project I!

write programs that:

- run on (multiple) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation

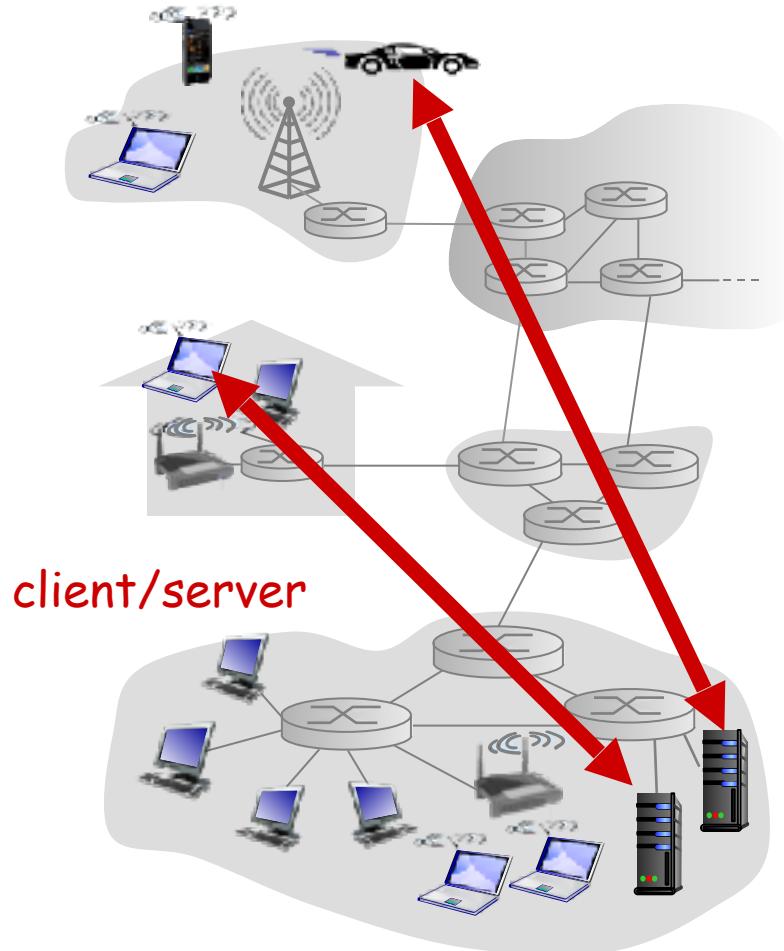


Application architectures

possible structure of applications:

- client-server
- peer-to-peer (P2P)

Client-server architecture



server:

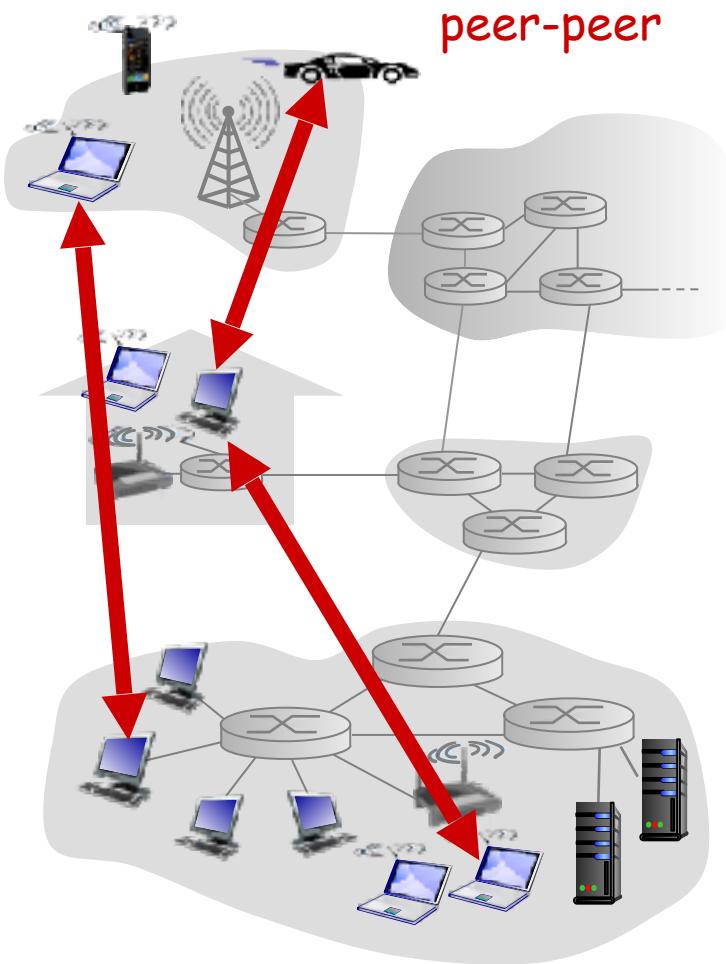
- always-on host
- permanent address
- data centers for scaling

clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management



Processes communicating

process: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**

clients, servers

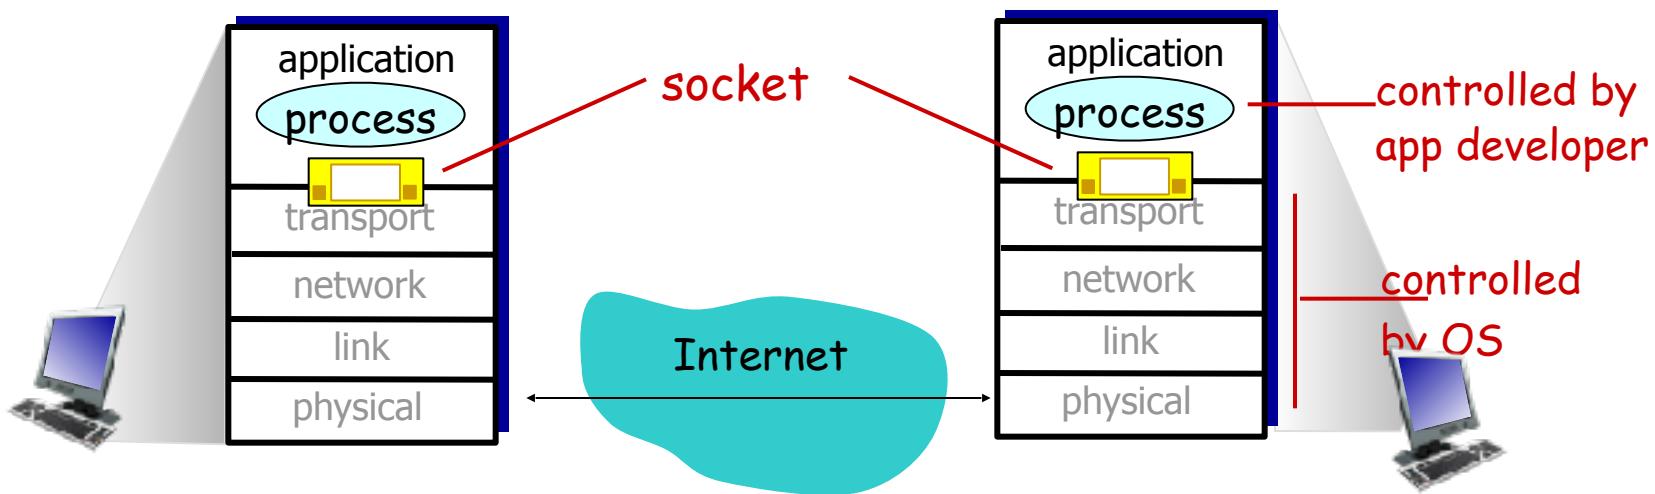
client process: process that initiates communication

server process: process that waits to be contacted

- § you'll write client & server processes in Project I
- § aside: applications with P2P architectures have client processes & server processes

Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
- upcoming class: socket programming in Python for Project I



Addressing processes

- to receive messages, process must have *identifier*
 - host device has unique 32-bit IP address
 - *Q:* does IP address of host on which process runs suffice for identifying the process?
- § *A:* no, many processes can be running on same host
- *identifier* includes both **IP address** and **port numbers** associated with process on host.
 - example port numbers:
 - HTTP server: 80
 - mail server: 25
 - to send HTTP message to `gaia.cs.umass.edu` web server:
 - **IP address:** 128.119.245.12
 - **port number:** 80
 - more shortly...

App-layer protocol defines

- types of messages exchanged,
 - e.g., request, response
- message syntax:
 - what fields in messages & how fields are delineated
- message semantics
 - meaning of information in fields
- rules for when and how processes send & respond to messages

open protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype

In coming classes, we will cover HTTP, SMTP/email, DNS, some P2P.

What properties might app need from transport service?

- ???

What transport service does an app need?

data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

throughput

- § some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- § other apps (“elastic apps”) make use of whatever throughput they get

security

- § encryption, data integrity, ...

Transport service requirements: common apps

application data loss

throughput

time sensitive

file transfer

e-mail

Web documents

real-time audio/video

stored audio/video

interactive games

text messaging

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100's msec
text messaging	no loss	elastic	yes and no

Internet transport protocols services

TCP service:

- *Provides:*
 - xxx?
- *Does not provide:*
 - xxx?

UDP service:

- *Provides:*
 - xxx?
- *Does not provide:*
 - xxx?

Internet transport protocols services

TCP service:

- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantee, security
- *connection-oriented*: setup required between client and server processes

UDP service:

- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliable or in order delivery, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

Recap: 4 dimensions of transport services

data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

throughput

- § some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- § other apps (“elastic apps”) make use of whatever throughput they get

timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

security

- § encryption, data integrity, ...

Internet apps: application, transport protocols

application	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	???
remote terminal access	Telnet [RFC 854]	???
Web	HTTP [RFC 2616]	???
file transfer	FTP [RFC 959]	???
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	???
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	???

Internet apps: application, transport protocols

application	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

Securing TCP

TCP & UDP

- no encryption
- if you send a cleartext password into socket, it traverses Internet in cleartext

SSL/TLS

- provides encrypted TCP connection
- data integrity
- end-point authentication
- SSL is older, deprecated

TLS is at app layer

- apps use SSL libraries, that “talk” to TCP
- Client/server negotiate use of TLS

TLS socket API

- if you send a cleartext password into socket, it traverses Internet encrypted

Chapter 2: outline

2.1 principles of applications
TCP vs UDP

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications
2.6 video streaming and
content distribution
networks

2.7 socket programming
with UDP and TCP

See *Shuyue's video, book*

**DO NOT SHARE
SLIDES AND CLASS MATERIALS
ON ONLINE SITES**

Uploading course materials to sites such as CourseHero, Chegg or Github is academic misconduct at Columbia (see [pg 10](#) of [Columbia guide](#)).

Day 6 1:10 video: Socket Programming



CSEE 4119
Computer Networks
Guest lecture by Shuyue Yu

 COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

Day 8: The Web/HTTP



CSEE 4119
Computer Networks
Ethan Katz-Bassett



IN THE CITY OF NEW YORK

Slides adapted from (and often identical to) slides from Kurose and Ross.

All material copyright 1996-2020

J.F Kurose and K.W. Ross, All Rights Reserved

Chapter 2: outline

- | | |
|--|---|
| 2.1 principles of applications
TCP vs UDP | 2.5 P2P applications |
| 2.2 Web and HTTP <i>Today</i> | 2.6 video streaming and
content distribution
networks |
| 2.3 electronic mail <ul style="list-style-type: none">• SMTP, POP3, IMAP | 2.7 socket programming
with UDP and TCP |
| 2.4 DNS | |

Web and HTTP

First, an overview...

- *web page* consists of *objects*
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,

www.someschool.edu/someDept/pic.gif

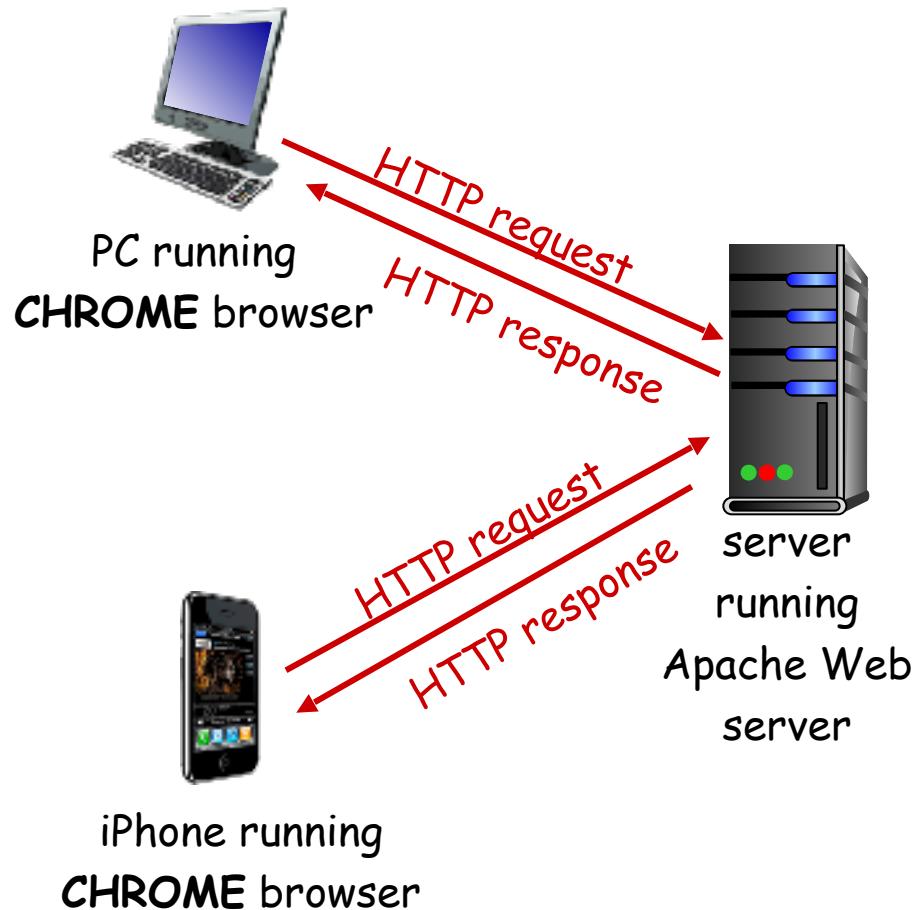
host name

path name

HTTP overview

HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - **client:** browser that requests, receives, (using HTTP protocol) and “displays” Web objects
 - **server:** Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

uses TCP:

- client initiates TCP connection (creates socket) to server, (usually) port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- From perspective of protocol, server maintains no information about past client requests

aside

protocols that maintain “state” are complex!

- § past history (state) must be maintained
- § if server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP connections

non-persistent HTTP

- at most one object sent over TCP connection
 - connection then closed
- downloading multiple objects required multiple connections

persistent HTTP

- multiple objects can be sent over single TCP connection between client, server

Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,
references to 10
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time
↓

Non-persistent HTTP (cont.)

time
↓

4. HTTP server closes TCP connection.

5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
6. Steps 1-5 repeated for each of 10 jpeg objects

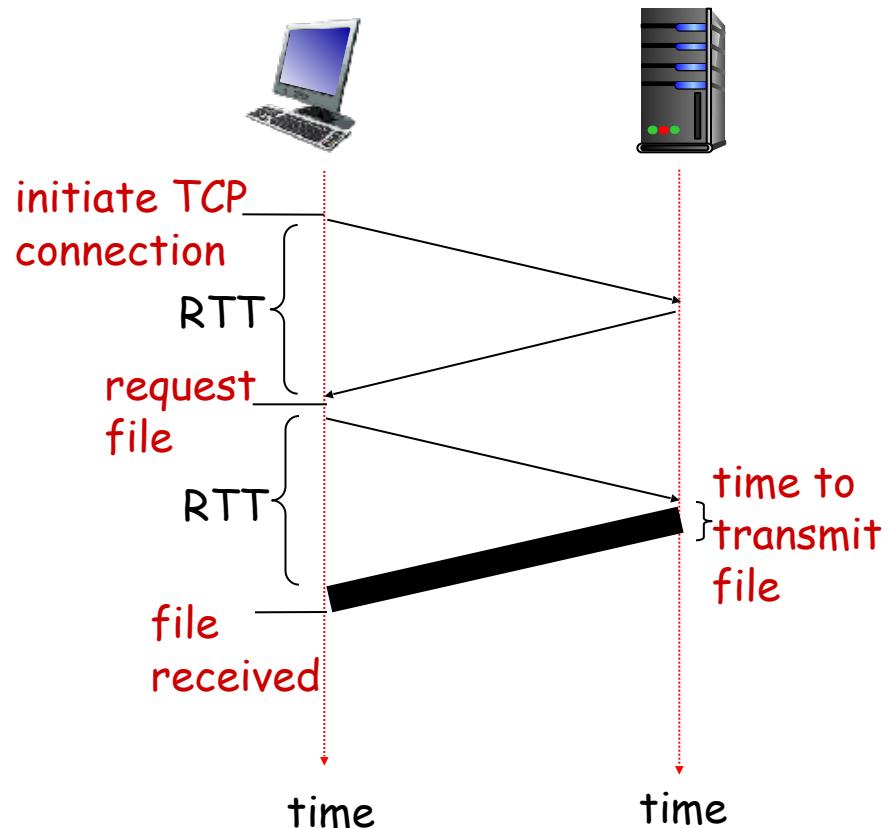
Non-persistent HTTP: response time

Round Trip Time (RTT) definition:

time for a small packet to travel from client to server and back

HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =
$$2\text{RTT} + \text{file transmission time}$$



Persistent HTTP

non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

persistent HTTP:

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

header
lines

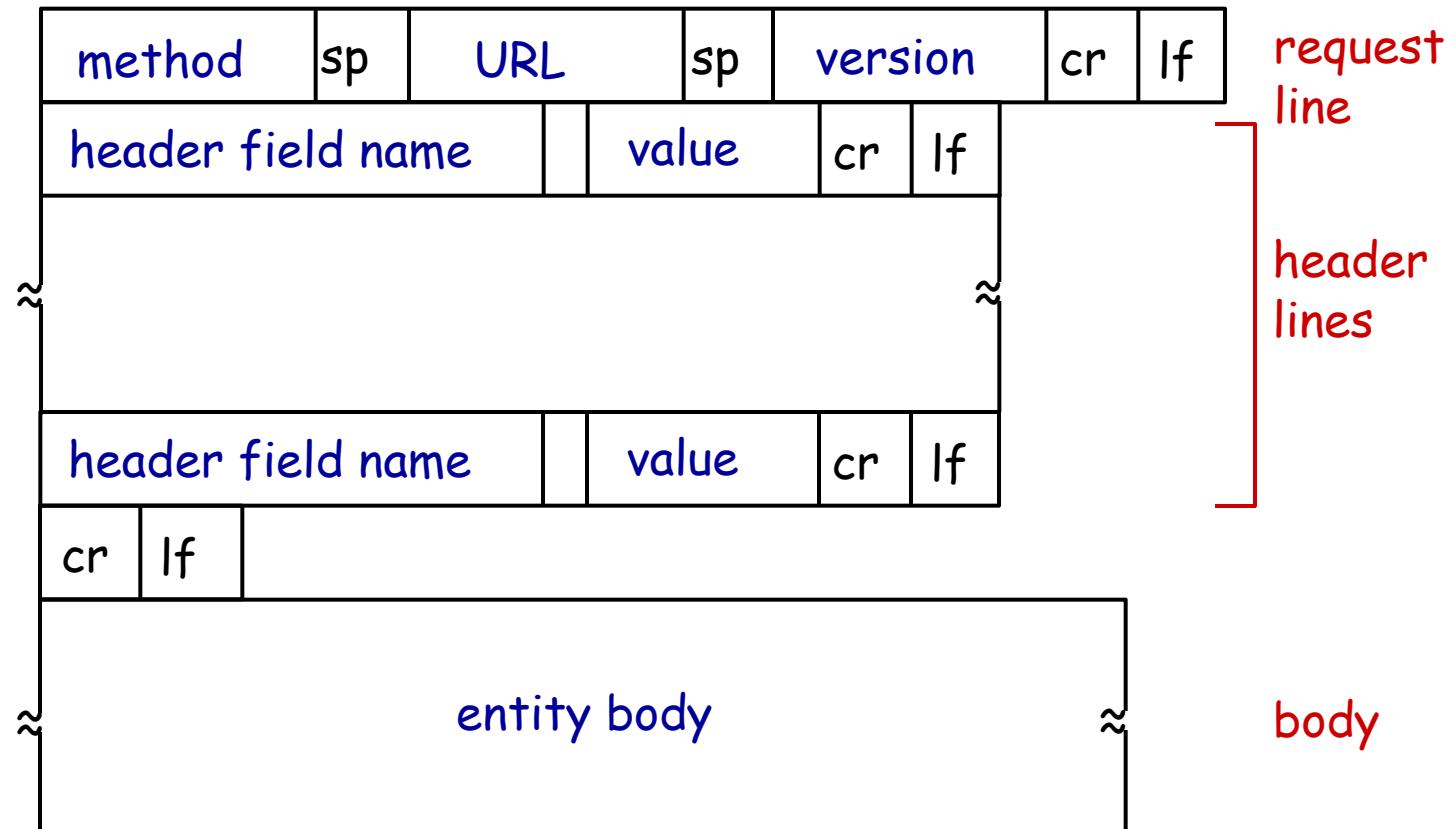
carriage return,
line feed at start
of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character
line-feed character

* Check out the online interactive exercises for more
examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

HTTP request message: general format



Uploading form input

POST method:

- web page often includes form input
- input is uploaded to server in entity body

URL method:

- uses GET method
- input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

Method types

HTTP/1.0:

- GET
- POST
- HEAD
 - asks server to leave requested object out of response

HTTP/1.1:

- GET, POST, HEAD
- PUT
 - uploads file in entity body to path specified in URL field
- DELETE
 - deletes file specified in the URL field

HTTP response message

status line

(protocol

status code

status phrase)

header
lines

data, e.g.,
requested
HTML file

```
HTTP/1.1 200 OK\r\nDate: Sun, 26 Sep 2010 20:09:20 GMT\r\nServer: Apache/2.0.52 (CentOS)\r\nLast-Modified: Tue, 30 Oct 2007 17:00:02  
GMT\r\nETag: "17dc6-a5c-bf716880"\r\nAccept-Ranges: bytes\r\nContent-Length: 2652\r\nKeep-Alive: timeout=10, max=100\r\nConnection: Keep-Alive\r\nContent-Type: text/html;  
charset=ISO-8859-1\r\n\r\ndata data data data data ...
```

* Check out the online interactive exercises for more
examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

HTTP response status codes

§ status code appears in 1st line in server-to-client response message.

§ some sample codes:

200 OK

- request succeeded, requested object later in this msg

301 Moved Permanently

- requested object moved, new location specified later in this msg
(Location:)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

`telnet gaia.cs.umass.edu 80`

opens TCP connection to port 80
(default HTTP server port)
at gaia.cs.umass.edu.
anything typed in will be sent
to port 80 at gaia.cs.umass.edu

2. type in a GET HTTP request:

`GET /kurose_ross/interactive/index.php HTTP/1.1`

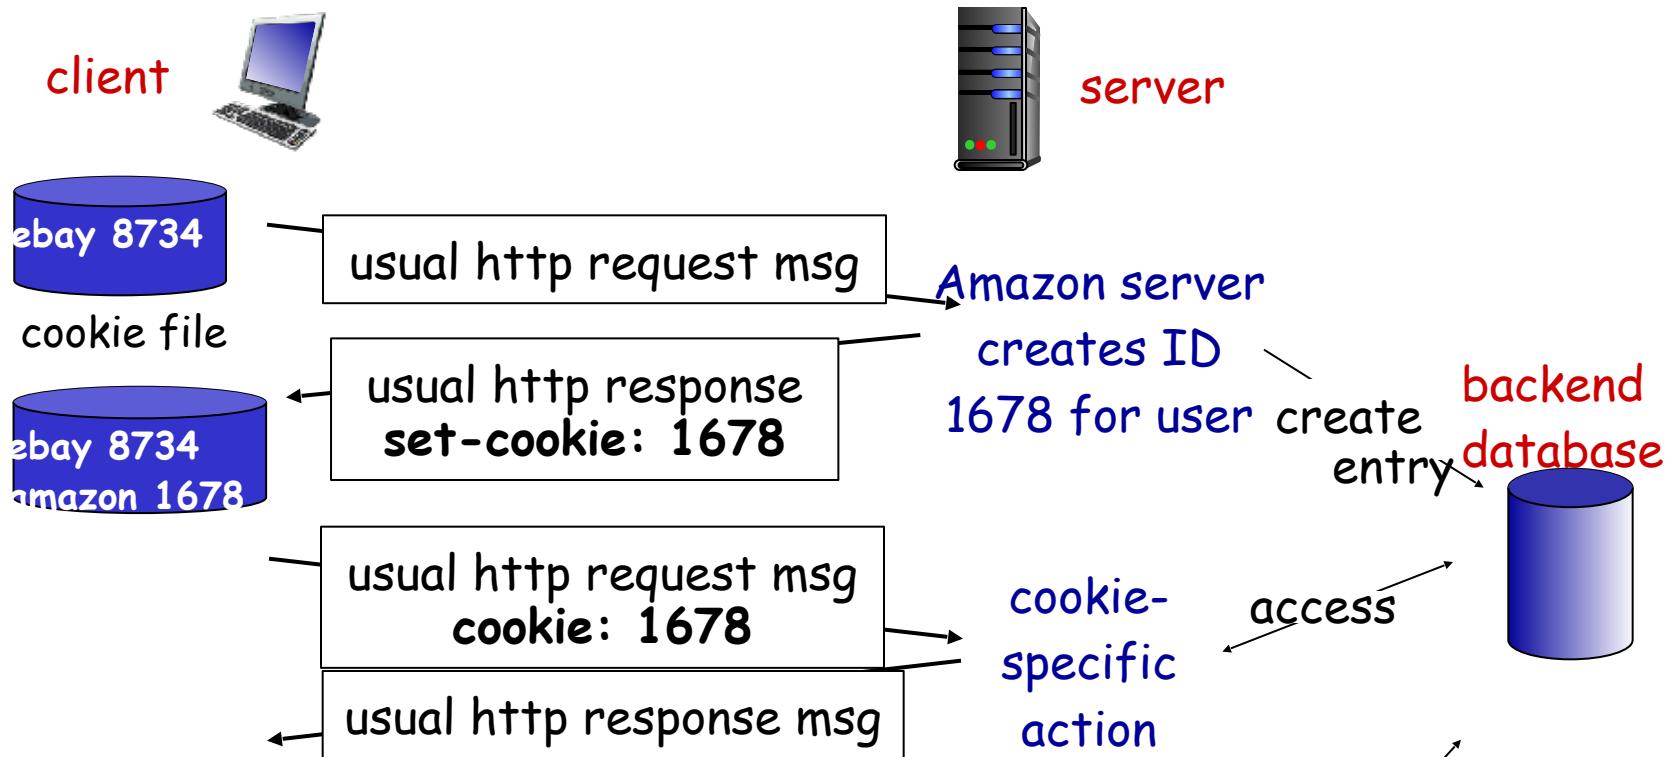
`Host: gaia.cs.umass.edu`

by typing this in (hit carriage
return twice), you send
this minimal (but complete)
GET request to HTTP server

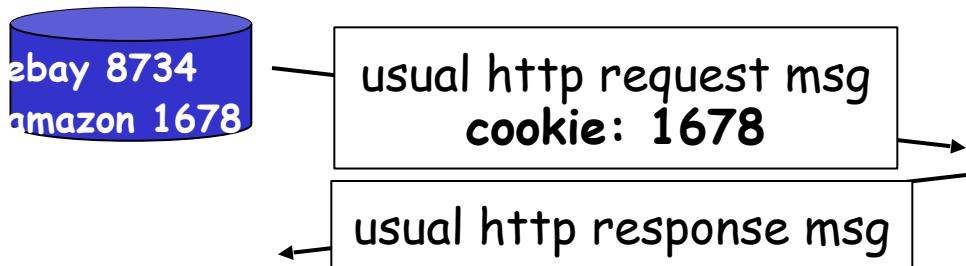
3. look at response message sent by HTTP server!

(or use Wireshark to look at captured HTTP request/response)

Cookies: keeping “state”



one week later:



User-server state: cookies

many Web sites use cookies

four components:

- 1) cookie header line of HTTP response message
- 2) cookie header line in next HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

example:

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

Cookies (continued)

what cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

how to keep “state”:

- § protocol endpoints: maintain state at sender/receiver over multiple transactions
- § cookies: http messages carry state

aside

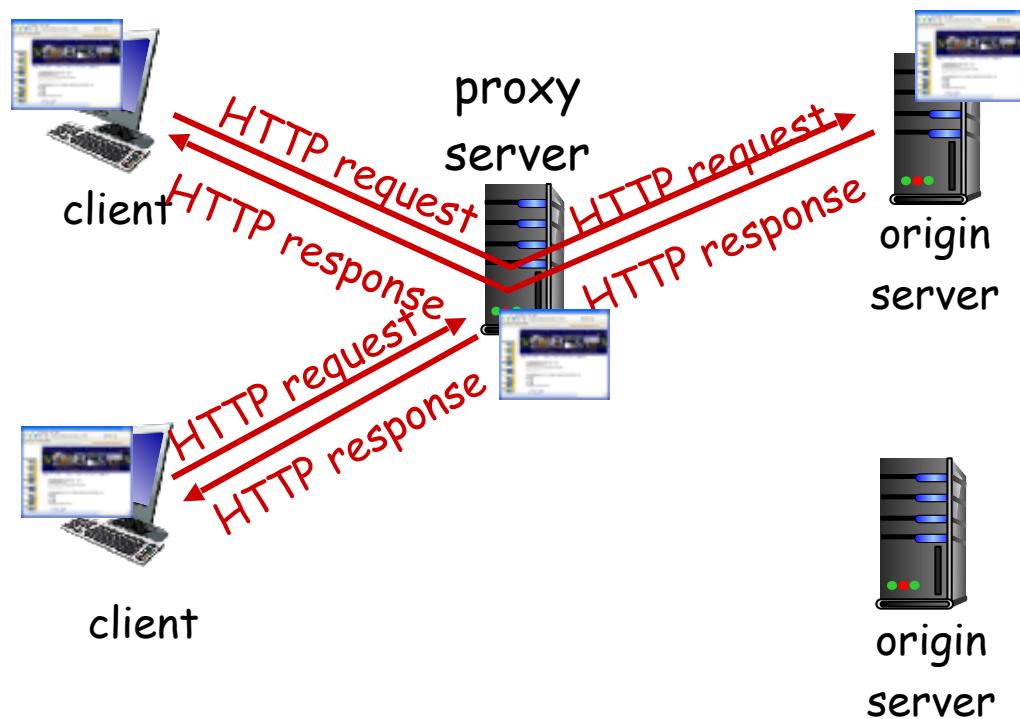
cookies and privacy:

- § cookies permit sites to learn a lot about you
- § you may supply name and e-mail to sites

Web caches (proxy server)

goal: satisfy request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- cache acts as both client and server
 - server for original requesting client
 - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)
 - also used by Content Delivery Networks (CDNs), which we will discuss next week
- Project 1: you will make a non-caching proxy

why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)

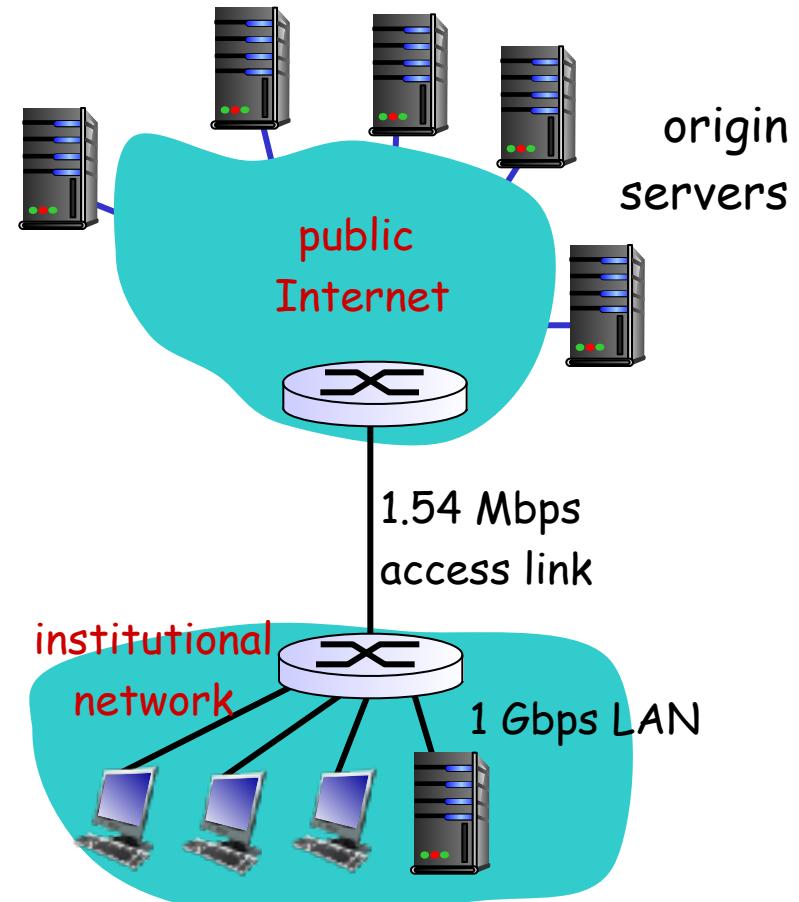
Caching example:

assumptions:

- § avg object size: 100K bits
- § avg request rate from browsers to origin servers: 15/sec
- § avg data rate to browsers: 1.50 Mbps
- § RTT from institutional router to any origin server: 2 sec
- § access link rate: 1.54 Mbps

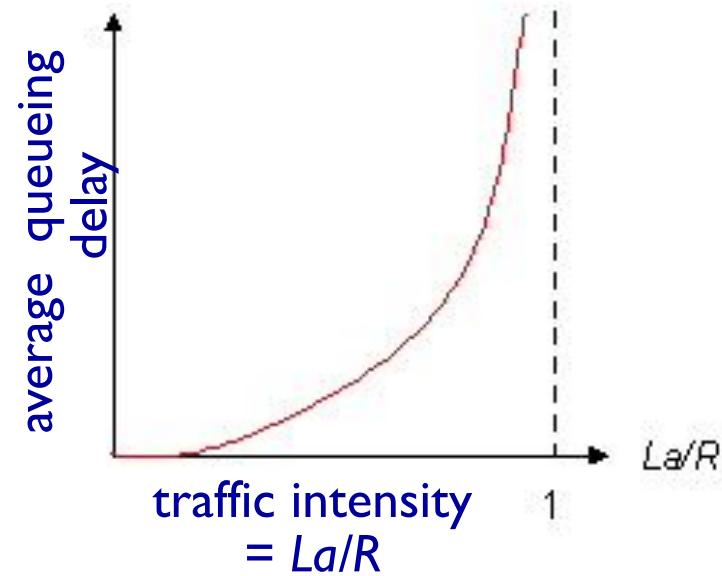
consequences:

- § LAN utilization: 0.15%
- § access link utilization = **97%**
- § total delay = Internet delay + access delay + LAN delay
= 2 sec + **minutes** + usecs



Reminder: Queueing delay

- R : link bandwidth (bps)
- L : packet length (bits)
- a : average packet arrival rate (packets per second)



- $La/R \sim 0$: avg. queueing delay small
- $La/R \rightarrow 1$: avg. queueing delay large
- $La/R > 1$: more “work” arriving than can be serviced, average delay infinite!



* Check online interactive animation on queuing and loss

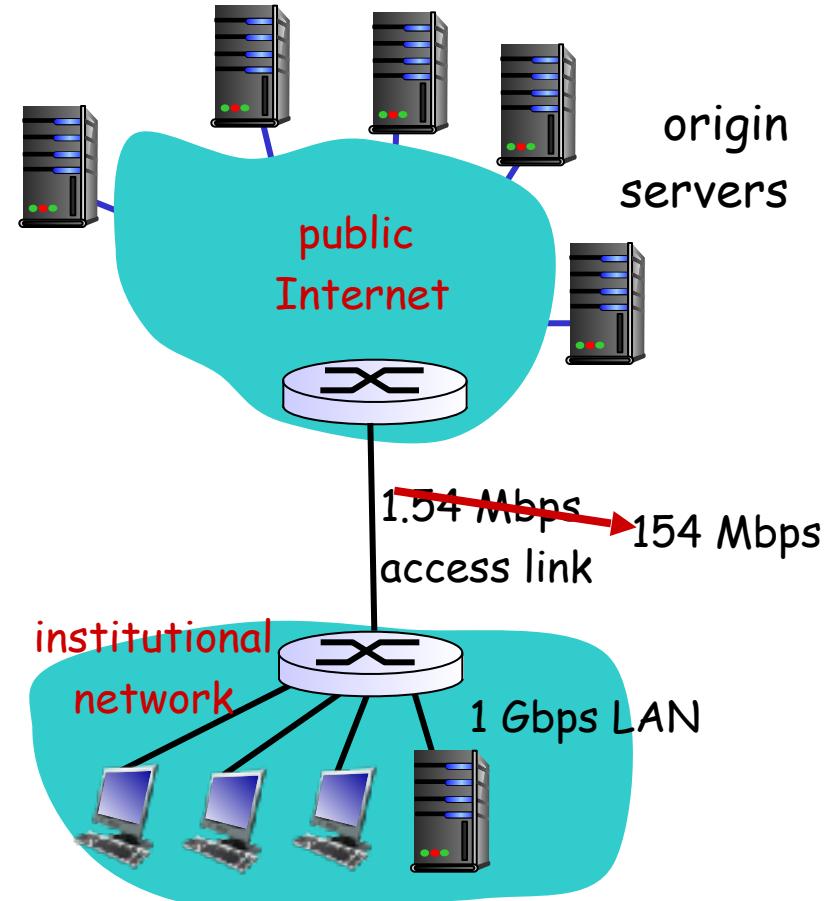
Caching example: fatter access link

assumptions:

- § avg object size: 100K bits
- § avg request rate from browsers to origin servers: 15/sec
- § avg data rate to browsers: 1.50 Mbps
- § RTT from institutional router to any origin server: 2 sec
- § access link rate: 1.54 Mbps

consequences:

- § LAN utilization: 0.15%
- § access link utilization = 97% → 0.97%
- § total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs
→ msec



Cost: increased access link speed (not cheap!)

Caching example: install local cache

assumptions:

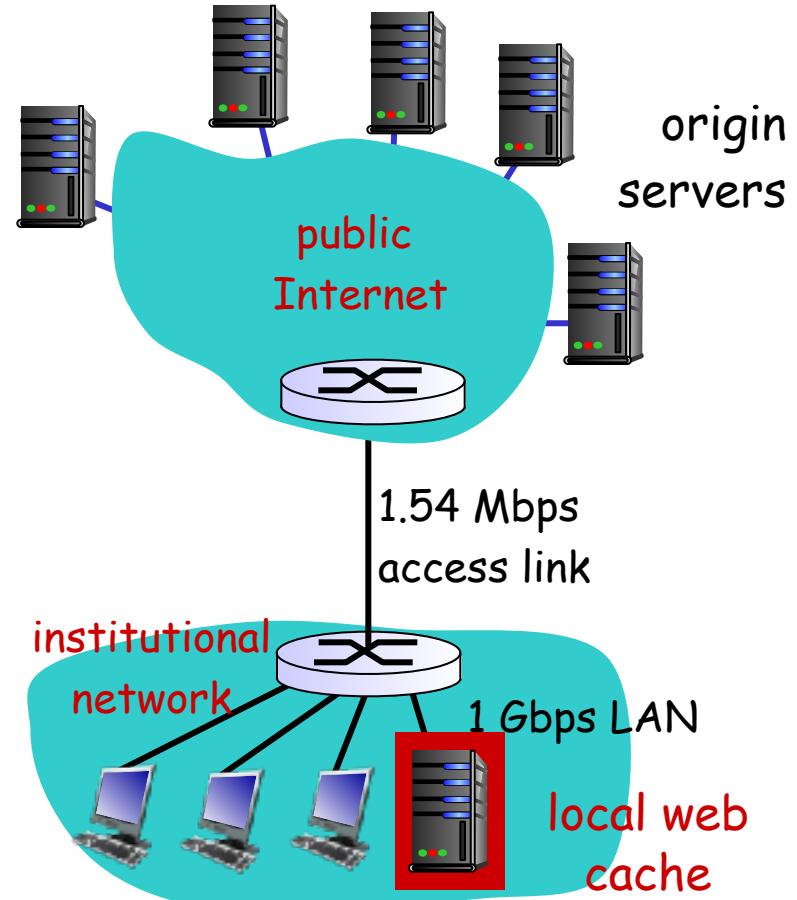
- § avg object size: 100K bits
- § avg request rate from browsers to origin servers: 15/sec
- § avg data rate to browsers: 1.50 Mbps
- § RTT from institutional router to any origin server: 2 sec
- § access link rate: 1.54 Mbps

consequences:

- § LAN utilization: 0.15%
- § access link utilization = ??
- § total delay = Internet delay? + access delay? + LAN delay
= ????

How to compute link utilization, delay?

Cost: web cache (cheap!)

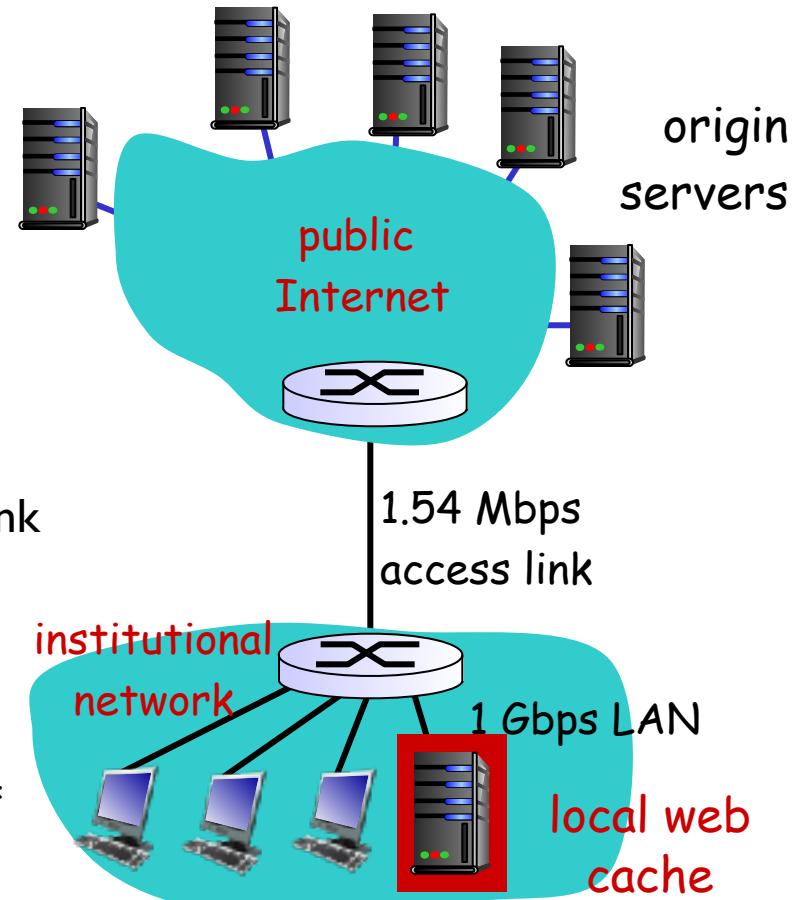


Caching example: install local cache

*Calculating access link utilization,
delay with cache:*

- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin

- § access link utilization:
 - § 60% of requests use access link
- § data rate to browsers over access link
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - § utilization = $0.9 / 1.54 = .58$
- § total delay
 - § $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - § $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
 - § less than with 154 Mbps link (and cheaper too!)



HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced multiple, pipelined GETs over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission
- even if server knows client is likely to request certain resources (e.g., images within a page client just requested), it has to wait for client to send them

HTTP/2

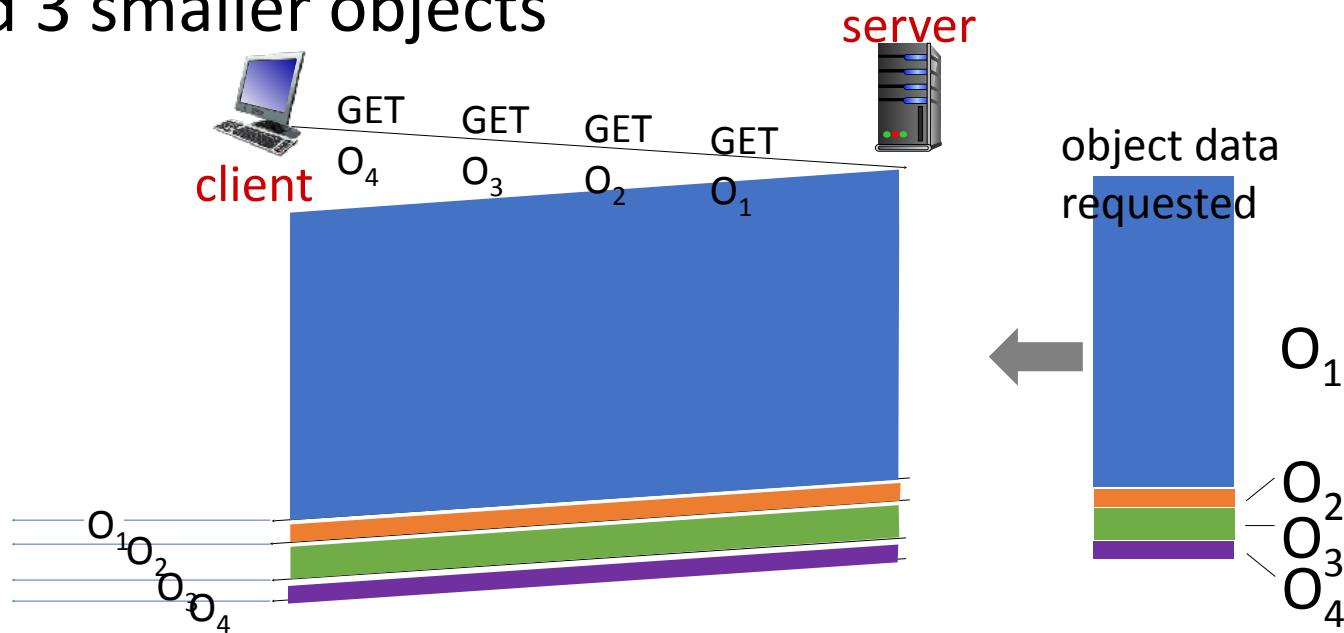
Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at server in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- interleave multiple requests/responses within same connection
 - divide objects into frames, schedule frames to mitigate HOL blocking

HTTP/2: mitigating HOL blocking

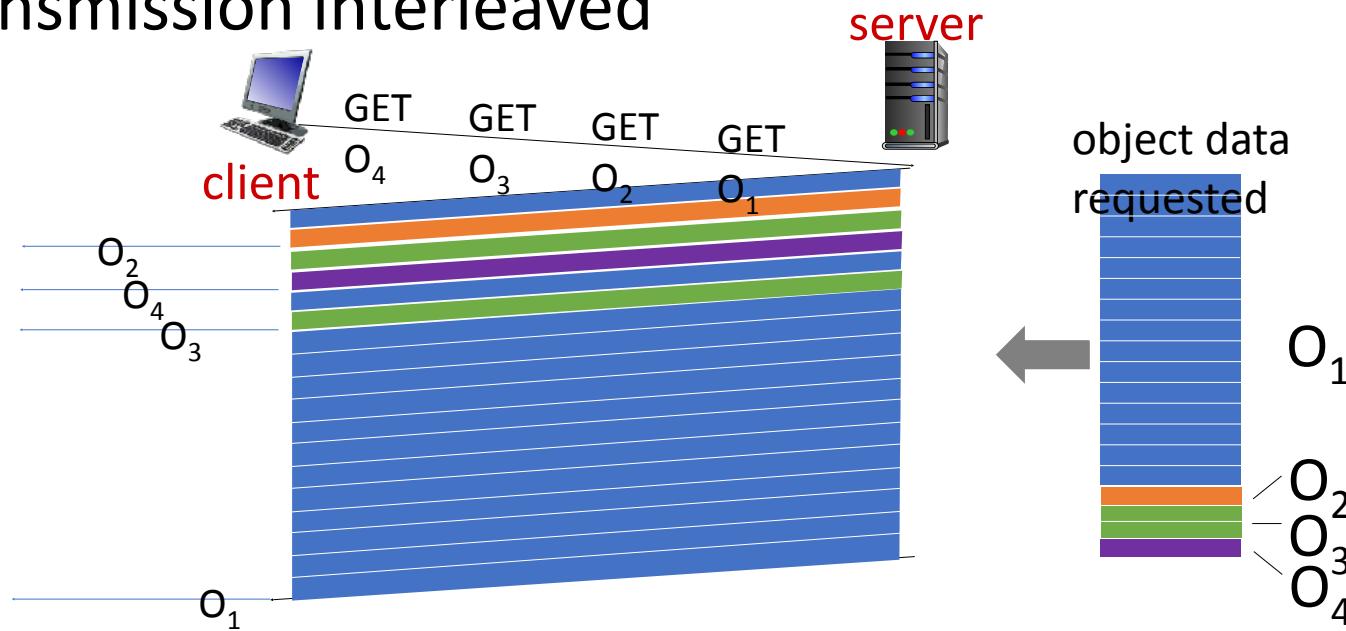
HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



*objects delivered in order requested: O₂, O₃, O₄
wait behind O₁*

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



O_2, O_3, O_4 delivered quickly, O_1 slightly delayed

HTTP/2 to HTTP/3

HTTP/2 over single TCP connection means:

- recovery from packet loss still stalls all object transmissions
 - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- no security over vanilla TCP connection
- **HTTP/3:** adds security, per object error- and congestion-control (more pipelining) over UDP

Chapter 2: outline

2.1 principles of applications
TCP vs UDP

2.2 Web and HTTP

2.3 electronic mail
• SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and
content distribution
networks

2.7 socket programming
with UDP and TCP

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

- caching

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

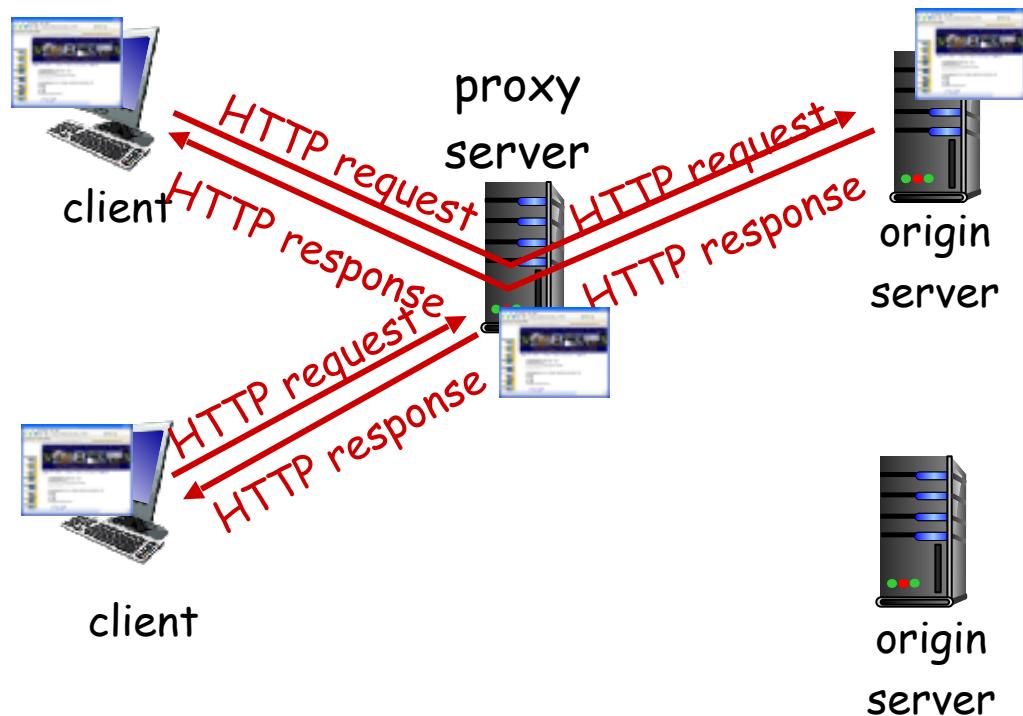
2.7 socket programming with UDP and TCP

Next

Recap: Web caches (proxy server)

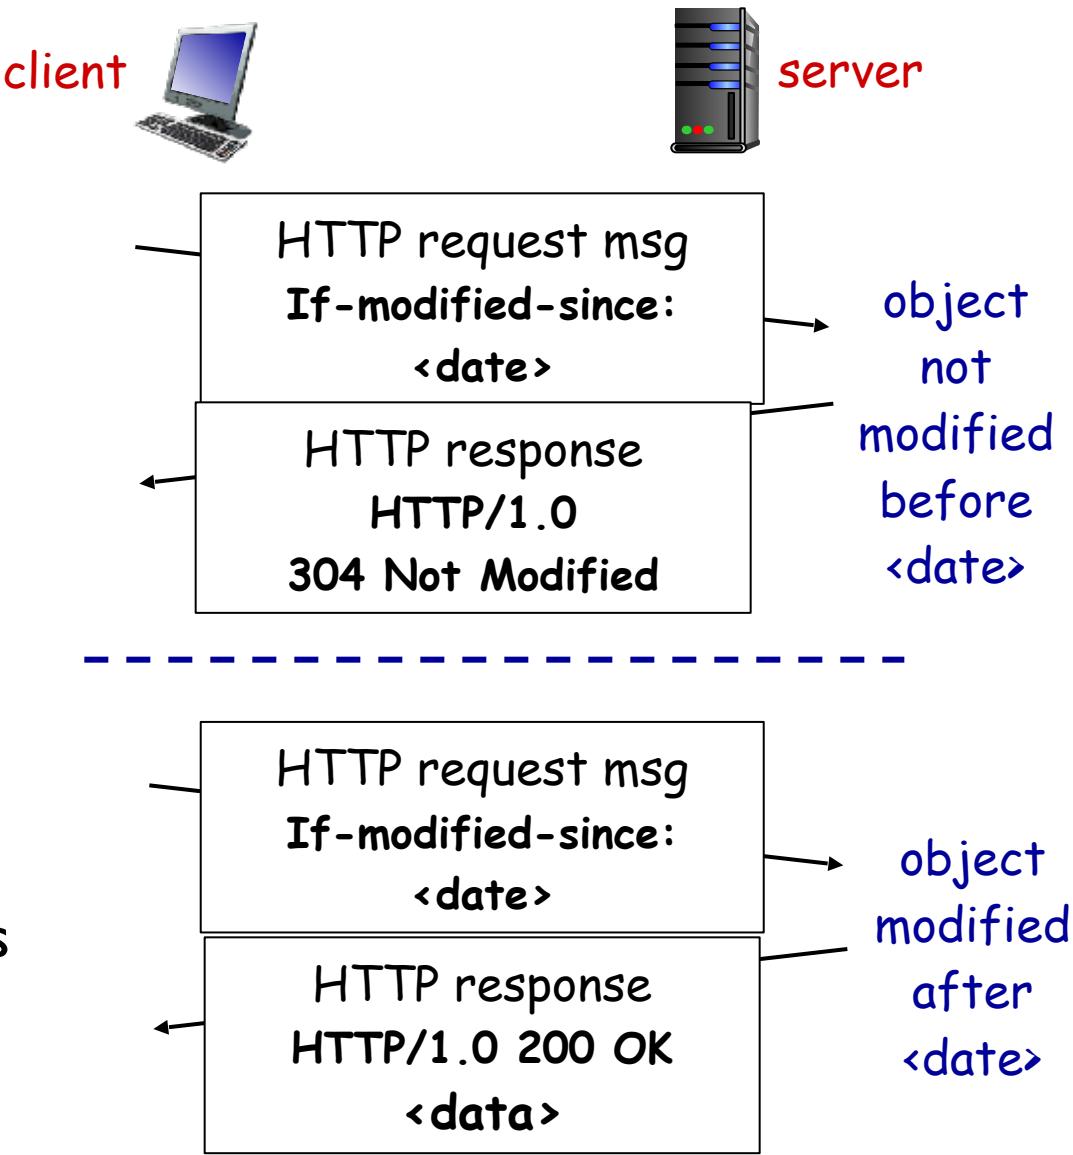
goal: satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
 - no object transmission delay
 - lower link utilization
- **cache:** specify date of cached copy in HTTP request
If-modified-since:
 `<date>`
- **server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

- caching

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

**DO NOT SHARE
SLIDES AND CLASS MATERIALS
ON ONLINE SITES**

Uploading course materials to sites such as CourseHero, Chegg or Github is academic misconduct at Columbia (see [pg 10](#) of [Columbia guide](#)).

Day 9: Email



CSEE 4119
Computer Networks
Ethan Katz-Bassett

 COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

Slides adapted from (and often identical to) slides from Kurose and Ross.

All material copyright 1996-2020

J.F.Kurose and K.W.Ross, All Rights Reserved

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

- caching

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS (as time allows)

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

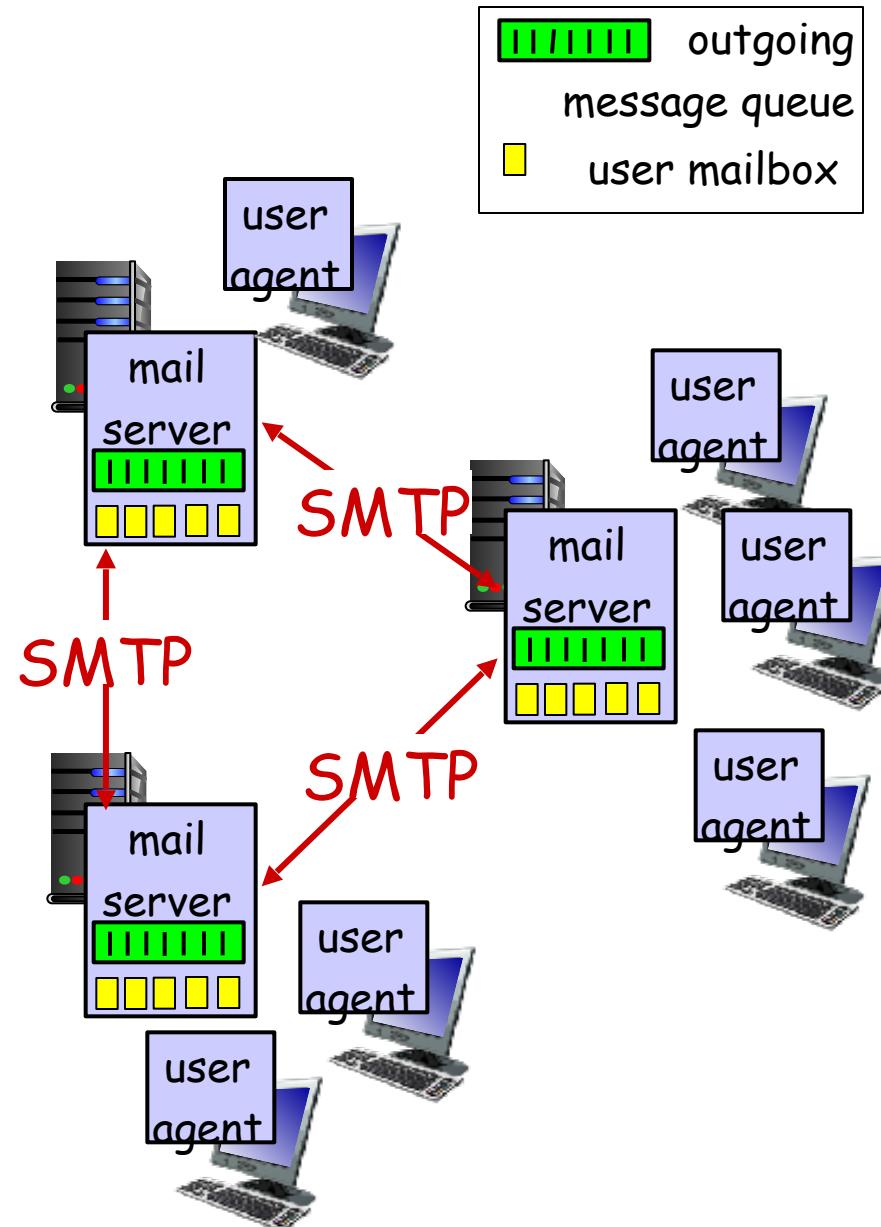
Electronic mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

User Agent

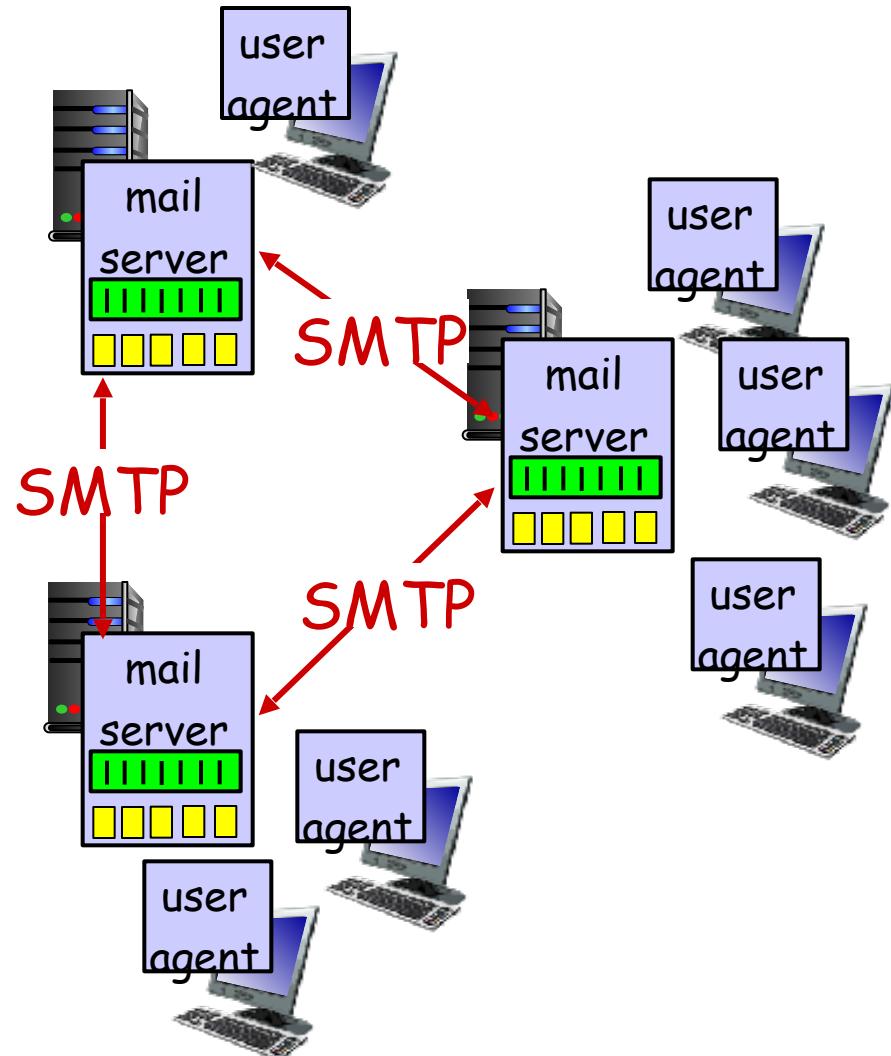
- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, Thunderbird, iPhone mail client
- outgoing, incoming messages stored on server



Electronic mail: mail servers

mail servers:

- *mailbox* contains incoming messages for user
- *message queue* of outgoing (to be sent) mail messages
- *SMTP protocol* between mail servers to send email messages
 - client: sending mail server
 - “server”: receiving mail server

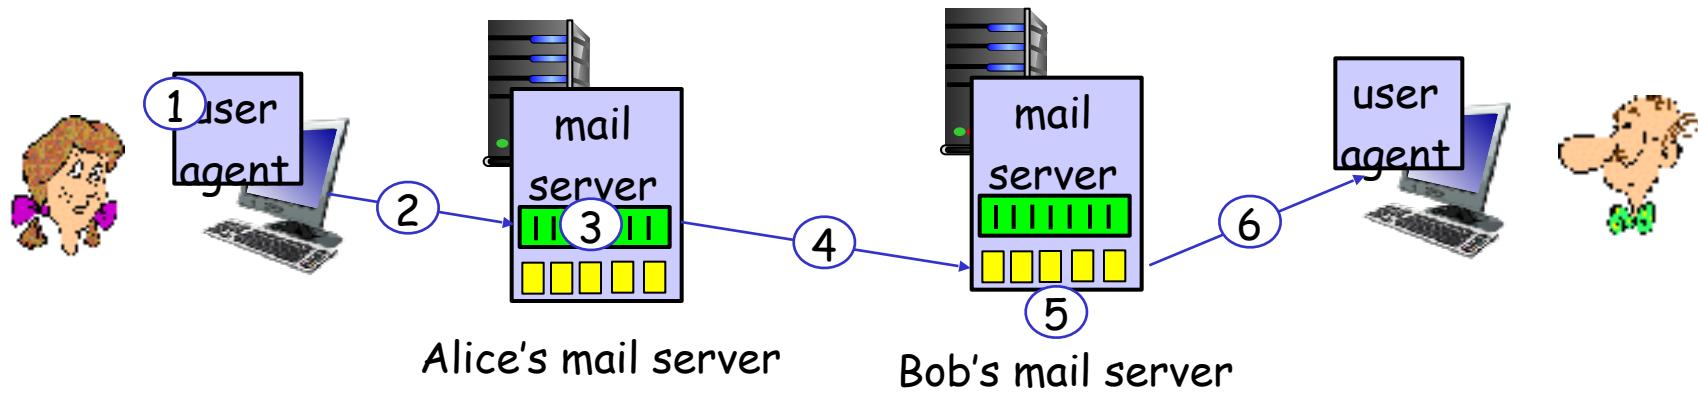


Electronic Mail: SMTP [RFC 2821]

- uses TCP to reliably transfer email message from client to server, port 25
- direct transfer: sending server to receiving server
- three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure
- command/response interaction (like HTTP)
 - commands: ASCII text
 - response: status code and phrase
- messages must be in 7-bit ASCII

Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server via SMTP; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

Try SMTP interaction for yourself:

- `telnet servername 25`
- see 220 reply from server
- enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

SMTP: final words

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF.CRLF to determine end of message

comparison with HTTP:

- HTTP: pull
- SMTP: push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in multipart message, encoded in ASCII

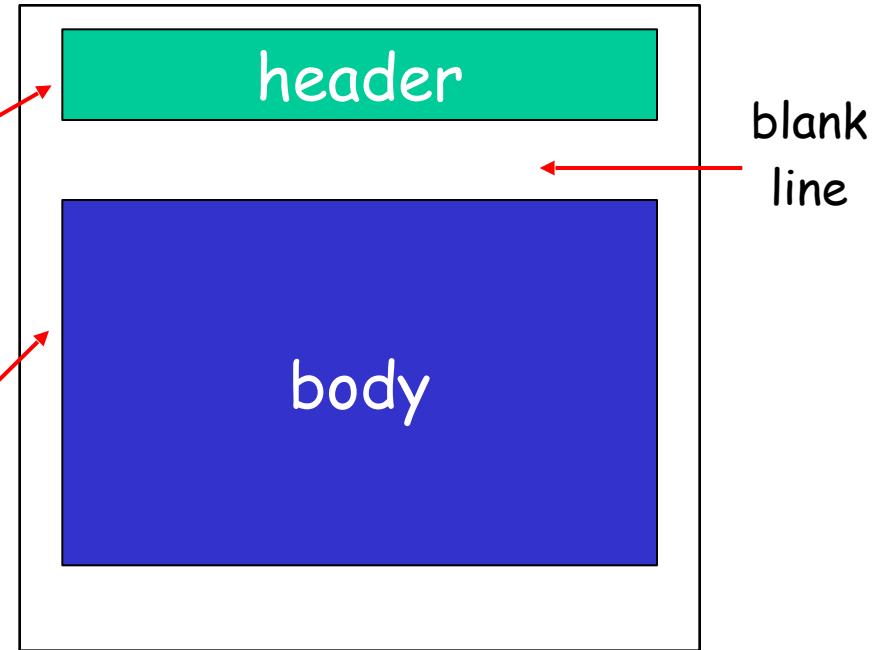
Mail message format

SMTP: protocol for exchanging
email messages

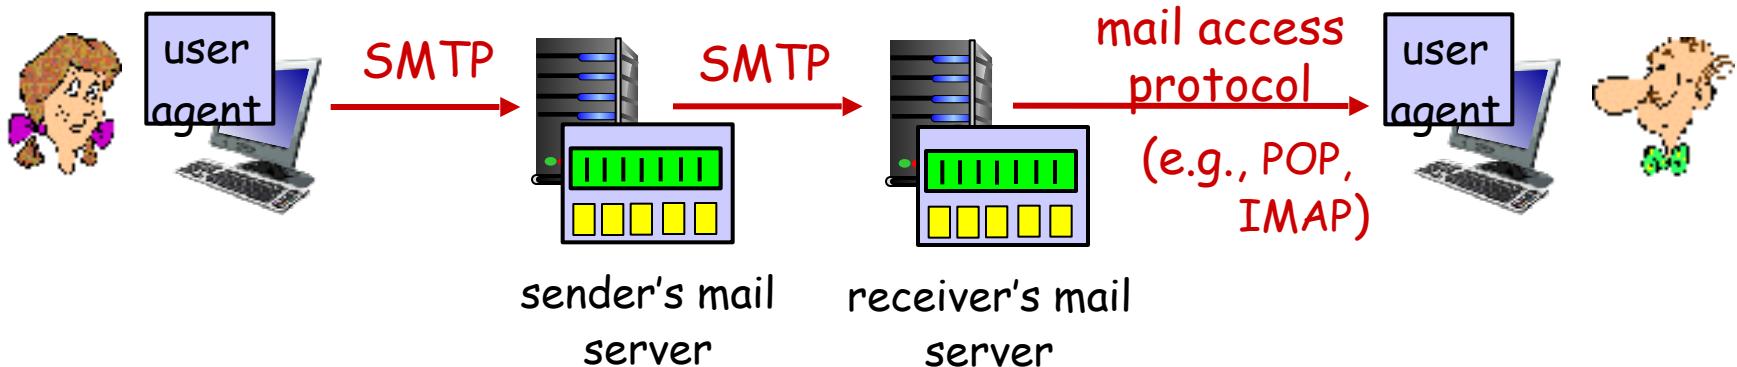
RFC 822: standard for text
message format:

- header lines, e.g.,
 - To:
 - From:
 - Subject:

*different from SMTP MAIL
FROM, RCPT TO:
commands!*
- Body: the “message”
 - ASCII characters only



Mail access protocols



- **SMTP:** delivery/storage to receiver's server
- mail access protocol: retrieval from server
 - **POP:** Post Office Protocol [RFC 1939]: authorization, download
 - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored messages on server
 - **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.

POP3 protocol

authorization phase

- client commands:
 - **user**: declare username
 - **pass**: password
- server responses
 - **+OK**
 - **-ERR**

transaction phase, client:

- **list**: list message numbers
- **retr**: retrieve message by number
- **dele**: delete
- **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

POP3 (more), IMAP, web mail

more about POP3

- previous example uses POP3 “download and delete” mode
 - Bob cannot re-read e-mail if he changes client
- POP3 “download-and-keep”: copies of messages on different clients
- POP3 is stateless across sessions

IMAP

- keeps all messages in one place: at server
- allows user to organize messages in folders
- keeps user state across sessions:
 - names of folders and mappings between message IDs and folder name

web mail

- send/receive mail via HTTP to mail server, which uses SMTP to other mail servers

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

- caching

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Next



**DO NOT SHARE
SLIDES AND CLASS MATERIALS
ON ONLINE SITES**

Course Hero

Uploading course materials to sites such as CourseHero, Chegg or Github is academic misconduct at Columbia (see [pg 10](#) of [Columbia guide](#)).

Day 10: DNS



CSEE 4119
Computer Networks
Ethan Katz-Bassett

 COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

Slides adapted from (and often identical to) slides from Kurose and Ross.

All material copyright 1996-2020

J.F Kurose and K.W. Ross, All Rights Reserved

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Next

DNS: domain name system

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans
 - other reasons?

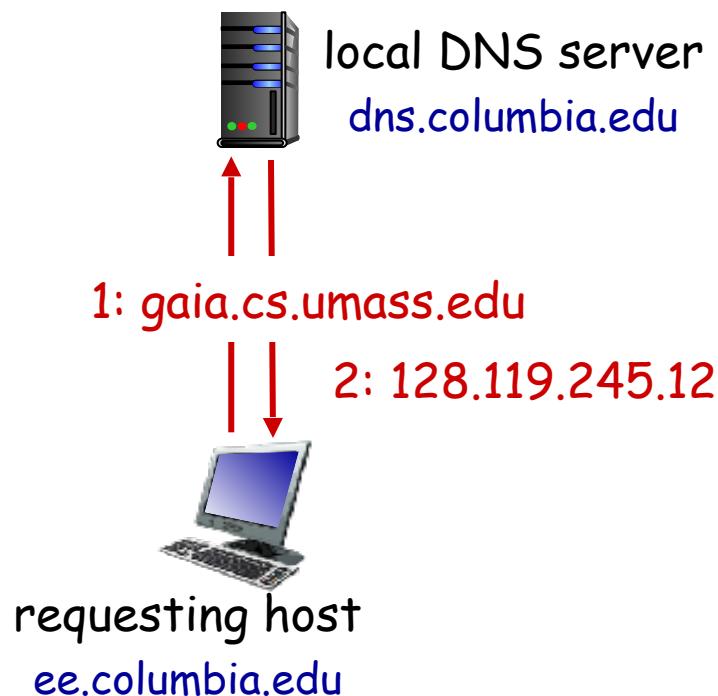
Q: how to map between IP address and name, and vice versa ?

Domain Name System:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)
 - note: core Internet function, implemented as application-layer protocol
 - used by many other applications and app-layer protocols
 - complexity at network’s “edge”

DNS name resolution: client view

- host at ee.columbia.edu wants IP address for gaia.cs.umass.edu
- asks local DNS server (usually hosted at local ISP)
 - UDP port 53
- from client's perspective, local DNS server responds with answer
 - rest of DNS system is a black box to client
 - we'll see how it works



DNS: services, structure

DNS services

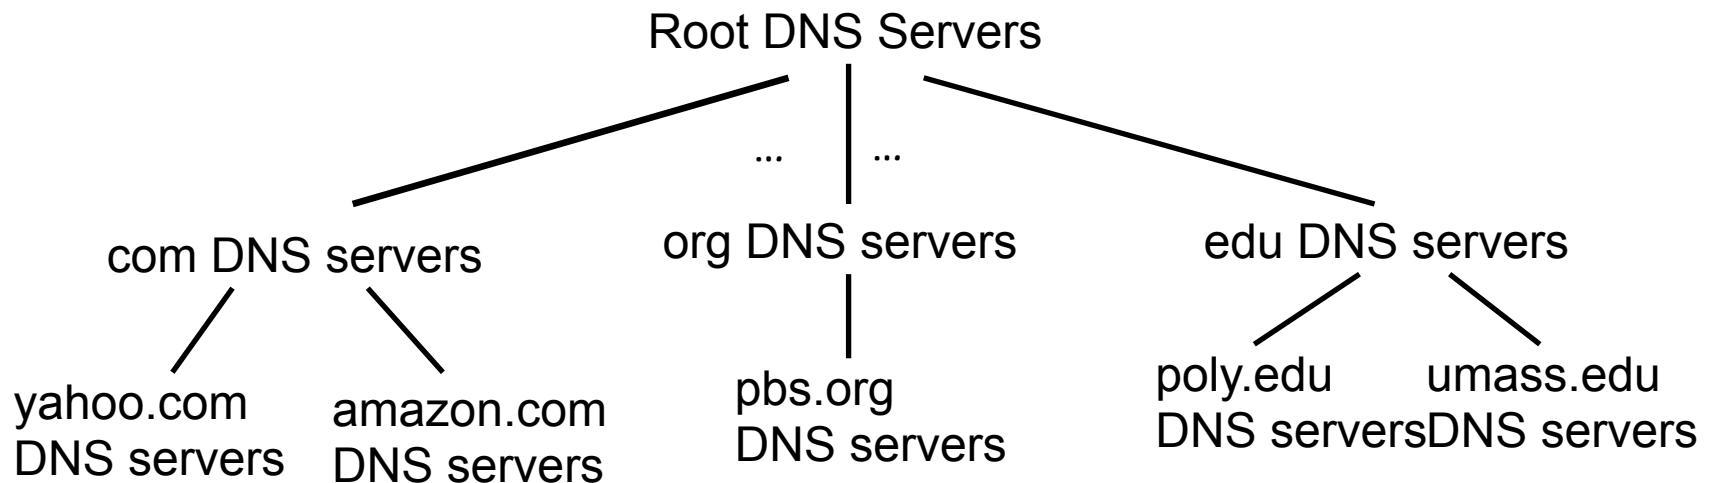
- hostname to IP address translation
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers:
many IP addresses
correspond to one
name

why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance/updates (e.g., new DNS names)

A: **doesn't scale!**

DNS: a distributed, hierarchical database



client wants IP for www.amazon.com; 1st approximation:

- client queries root server to find .com DNS server (IP addresses)
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

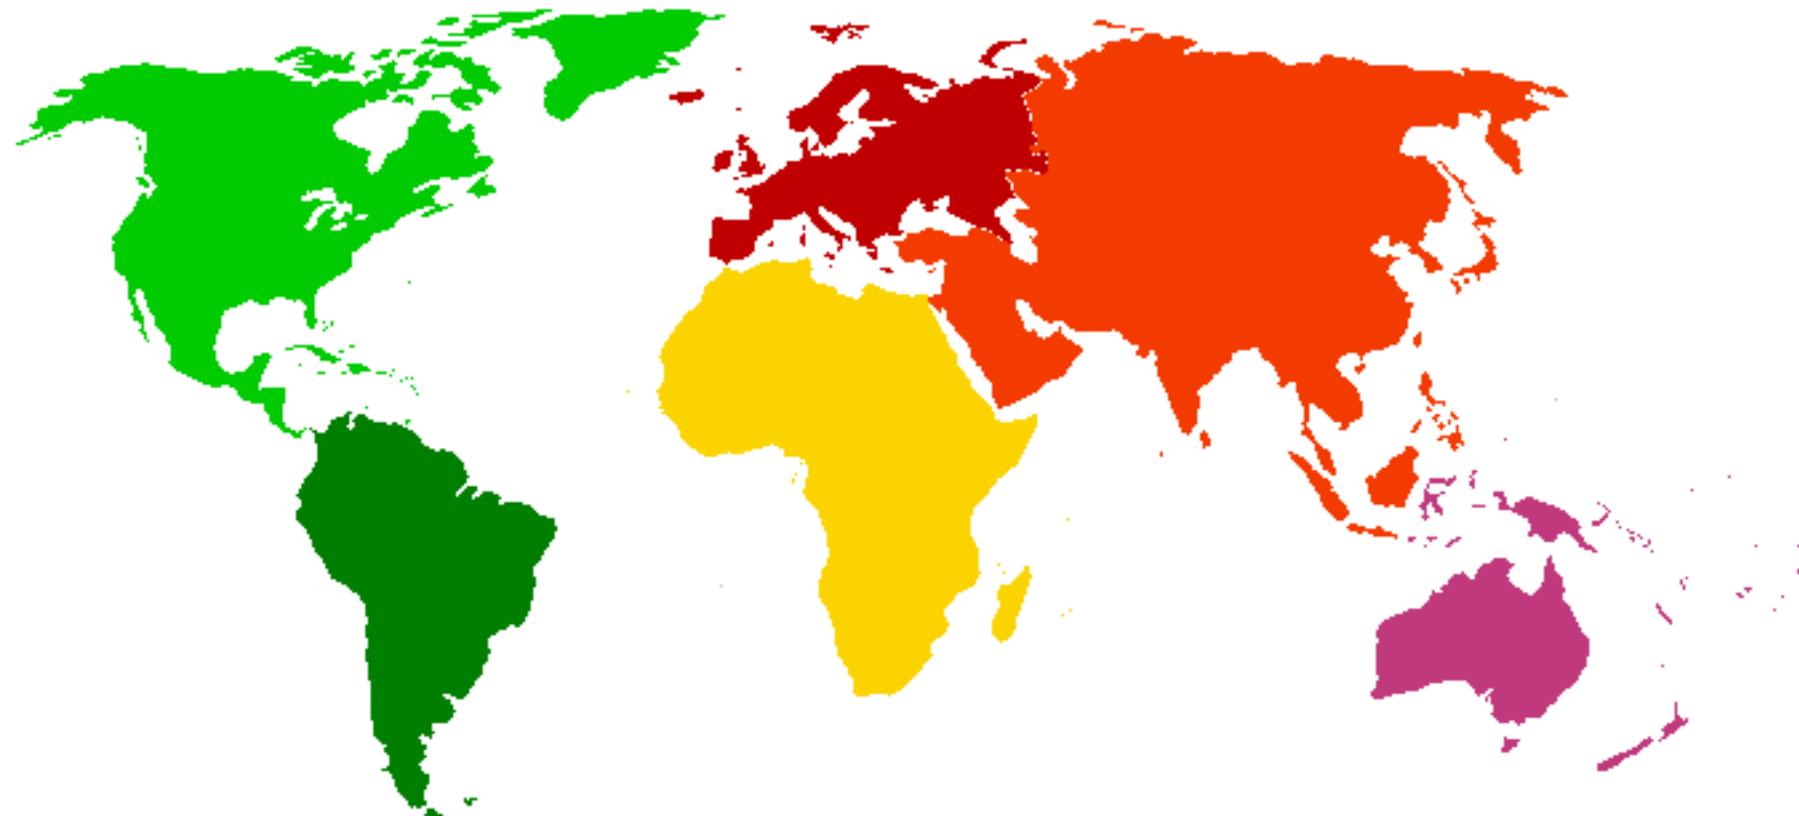
DNS: root name servers

- contacted by local name server that cannot resolve name
 - can contact any, regardless of what query
 - root server IP addresses loaded into your ISP's resolver

13 logical root name “servers” worldwide, labeled a...m

Where are they? To start answering:

What latencies did we see for c, e, and j?

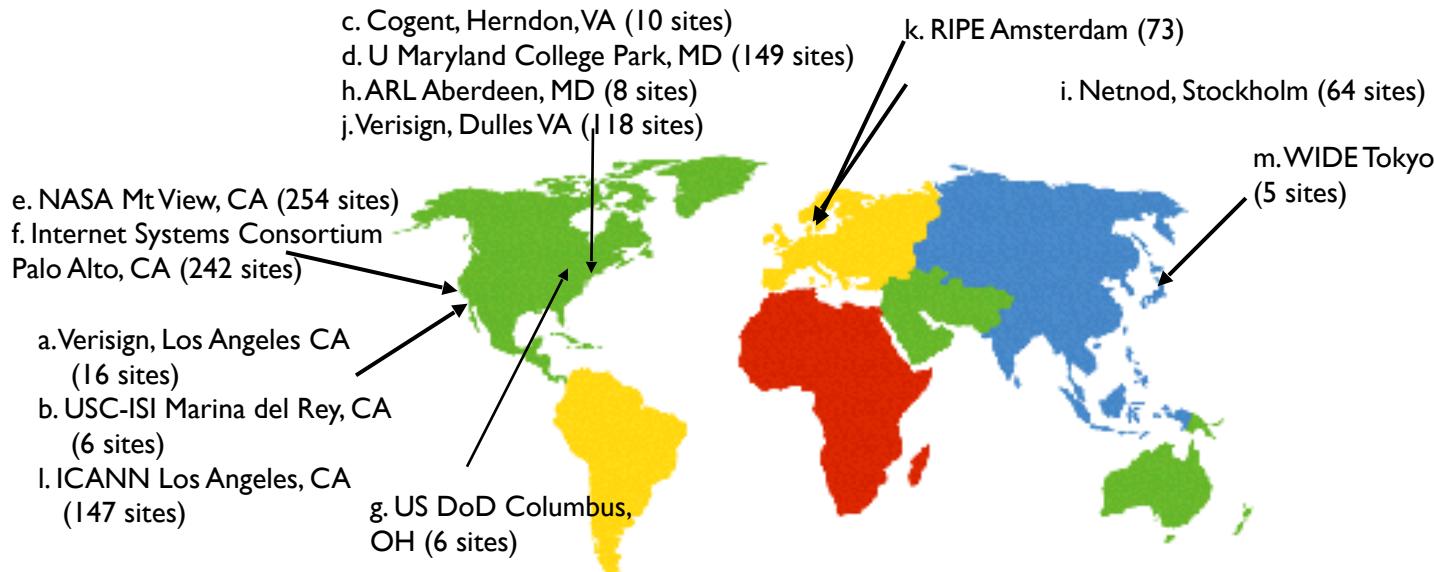


DNS: root name servers

- contacted by local name server that cannot resolve name
 - can contact any, regardless of what query
 - root server IP addresses loaded into your ISP's resolver

13 logical root name
“servers” worldwide

- each “server” replicated many times
- numbers updated 10/2020



TLD, authoritative servers

top-level domain (TLD) servers:

- com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp each have
- VeriSign maintains servers for .com TLD
- Educause for .edu TLD

authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- Every organization with publicly accessible hosts (such as Web servers and mail servers) must provide publicly accessible DNS records to map those hosts to IP addresses
- can be maintained by organization or service provider

Local DNS name server

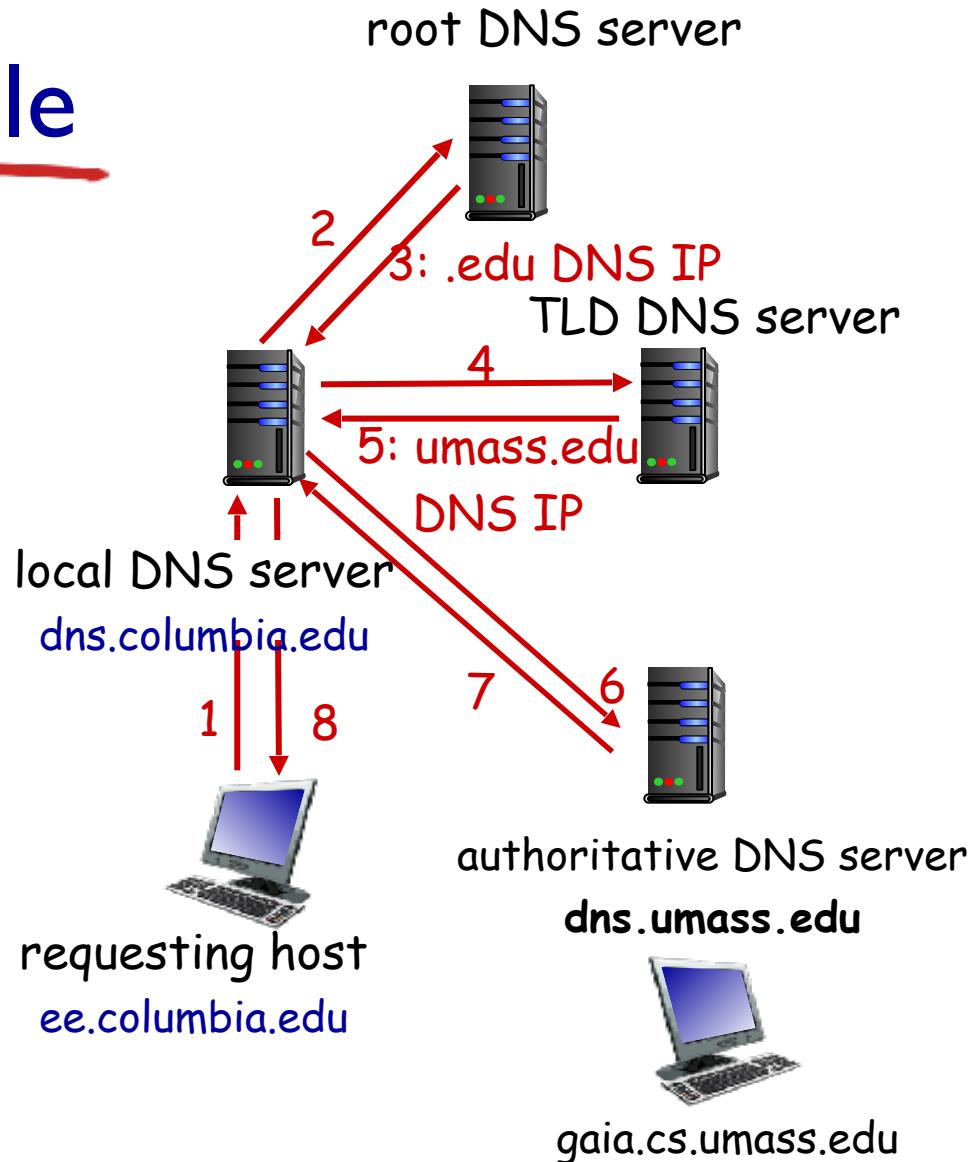
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
 - also called “default name server”
 - also public DNS servers:
Google public DNS (8.8.8.8), OpenDNS
- when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

DNS name resolution example

- host at ee.columbia.edu wants IP address for gaia.cs.umass.edu

iterated query:

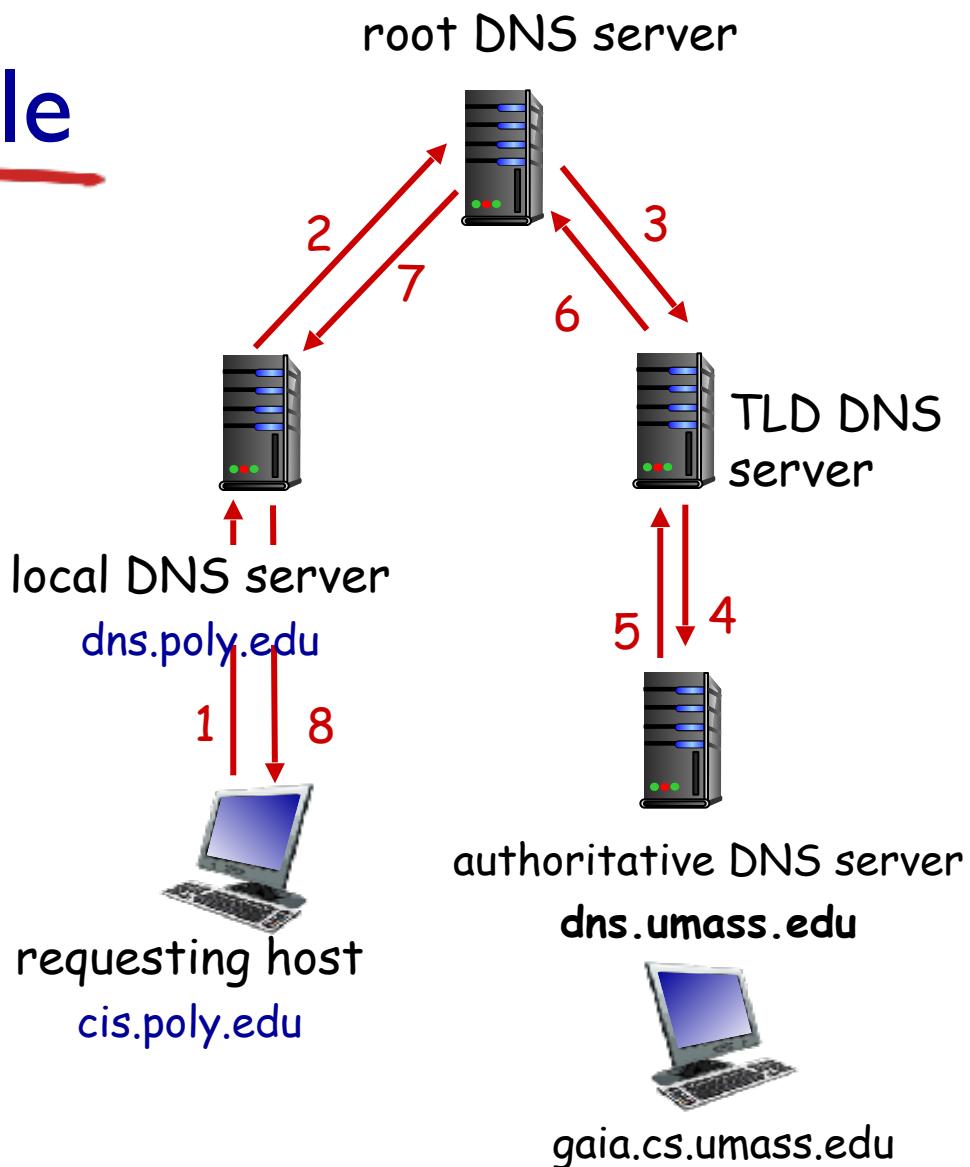
- § contacted server replies with name of server to contact
- § “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- § puts burden of name resolution on contacted name server
- § heavy load at upper levels of hierarchy?



DNS: caching, updating records

- once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
 - thus root name servers not often visited
- cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard
 - RFC 2136

DNS records

DNS: distributed database storing resource records (**RR**)

RR format: `(name, value, type, ttl)`

type=A

- § **name** is hostname
- § **value** is IP address

type=NS

- **name** is domain (e.g.,
`foo.com`)
- **value** is hostname of
authoritative name server
for this domain

type=CNAME

- § **name** is alias name for some
“canonical” (the real) name
- § `www.ibm.com` is really
`servereast.backup2.ibm.com`
- § **value** is canonical name

type=MX

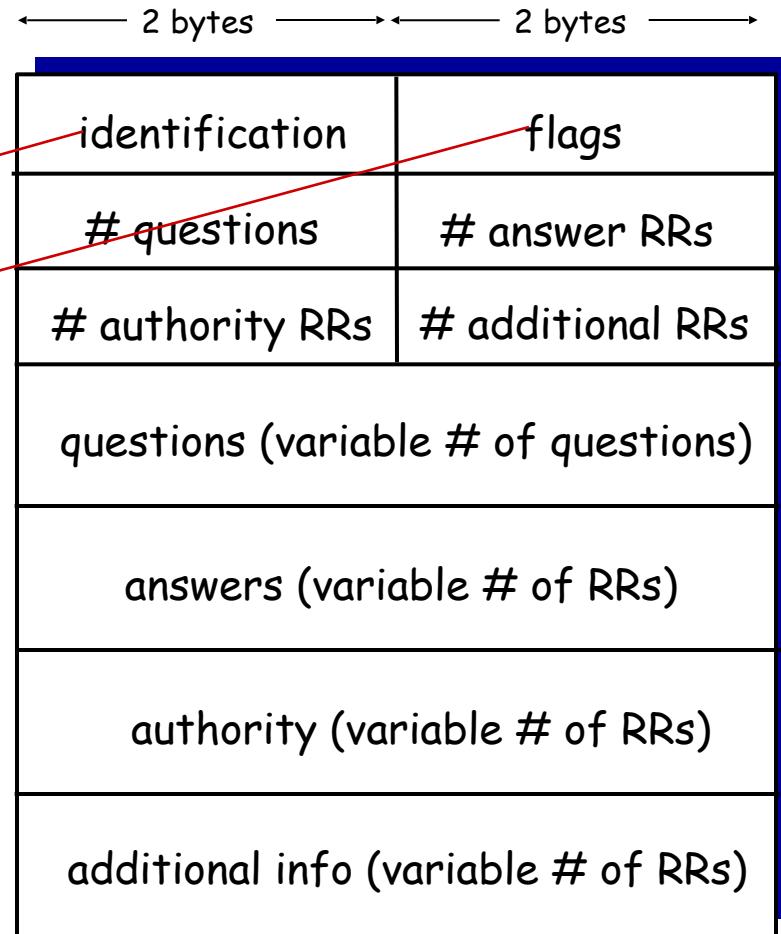
- § **value** is name of mailserver
associated with **name**

DNS protocol, messages

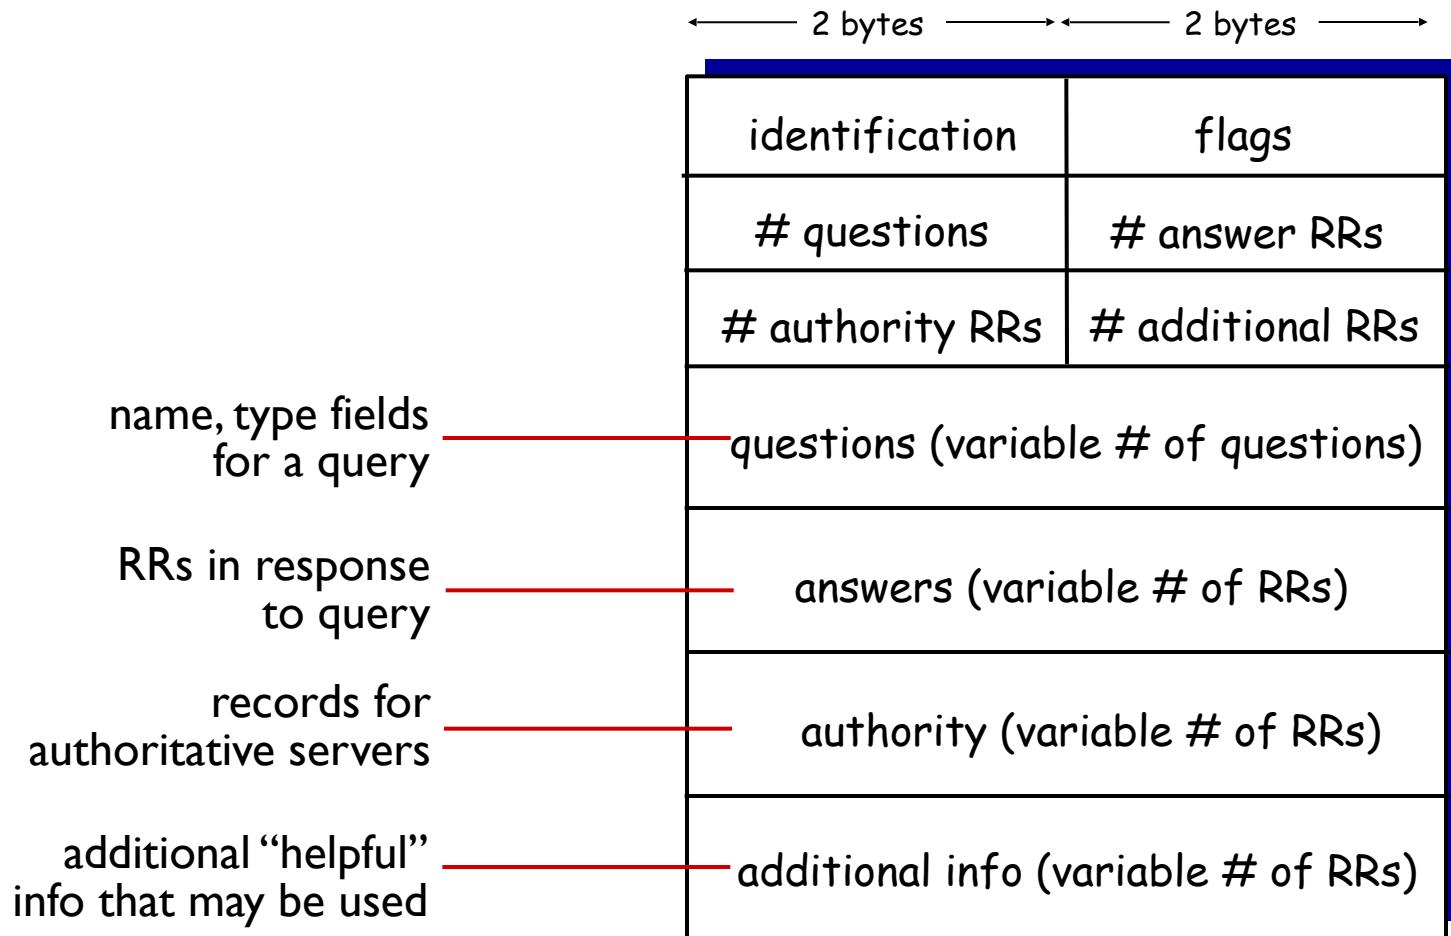
- *query* and *reply* messages, both with same *message format*

message header

- § identification: 16 bit # for query, reply to query uses same #
- § flags:
 - § ~~query or reply~~
 - § recursion desired
 - § recursion available
 - § reply is authoritative



DNS protocol, messages



Inserting records into DNS

- example: new startup “Network Utopia”
- register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server (primary and secondary)
 - registrar inserts two RRs into .com TLD server:
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server type A record for www.networkutopia.com; type MX record for networkutopia.com

Attacking DNS

DDoS attacks

- bombard root servers with traffic
 - some attacks last year!!
 - traffic filtering
 - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
 - potentially more dangerous

redirect attacks

- man-in-middle
 - Intercept queries
- DNS poisoning
 - Send bogus replies to DNS server, which caches

exploit DNS for DDoS

- send queries with spoofed source address: target IP
- requires amplification

**DO NOT SHARE
SLIDES AND CLASS MATERIALS
ON ONLINE SITES**

Course Hero

Uploading course materials to sites such as CourseHero, Chegg or Github is academic misconduct at Columbia (see [pg 10](#) of [Columbia guide](#)).

Day 11: Distributing Video and Other Content



CSEE 4119
Computer Networks
Ethan Katz-Bassett

 COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

Slides adapted from (and often identical to) slides from Kurose and Ross.

All material copyright 1996-2020

J.F.Kurose and K.W.Ross, All Rights Reserved

Chapter 2: outline

Today's class

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

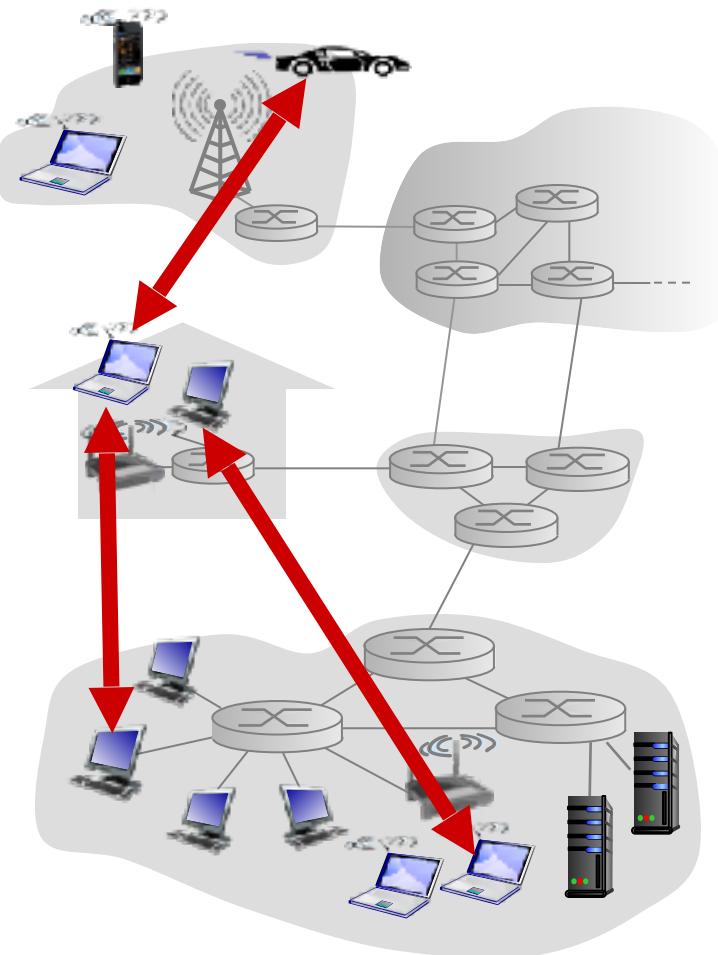
2.7 socket programming with UDP and TCP

Pure P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:

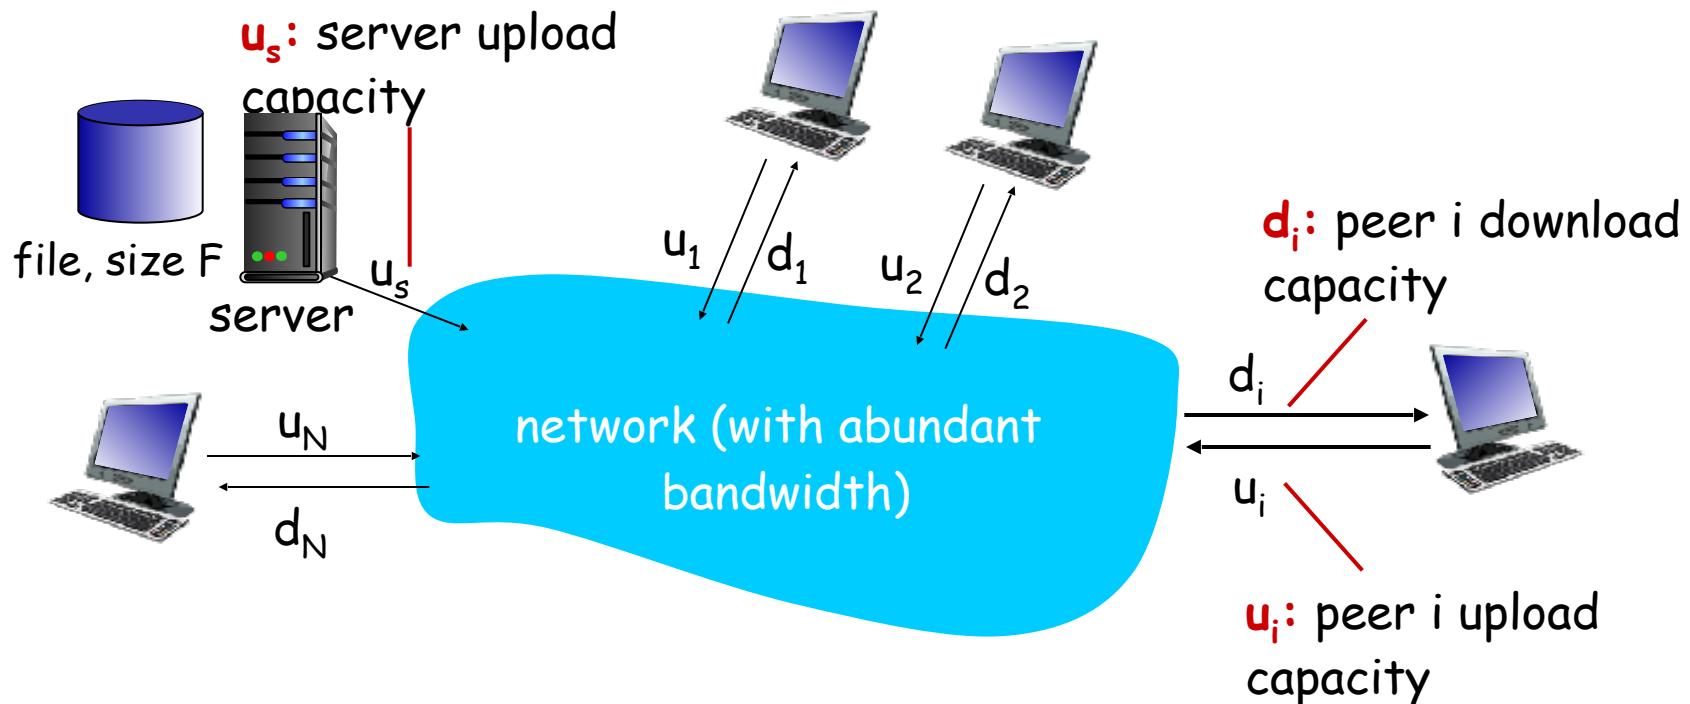
- file distribution (BitTorrent)
- Streaming (PPTV)
- VoIP (Skype...maybe)



File distribution: client-server vs P2P

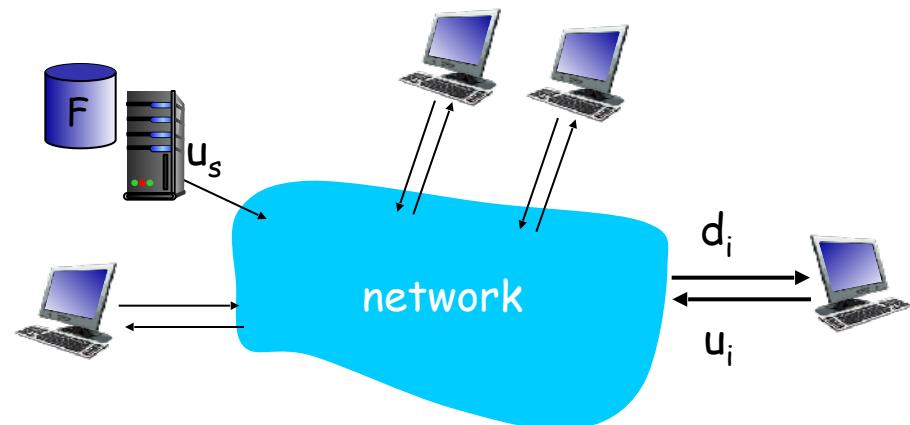
Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:
 - time to send one copy: F/u_s
 - time to send N copies: NF/u_s



- § **client:** each client must download file copy
 - d_{min} = min client download rate
 - max client download time: F/d_{min}

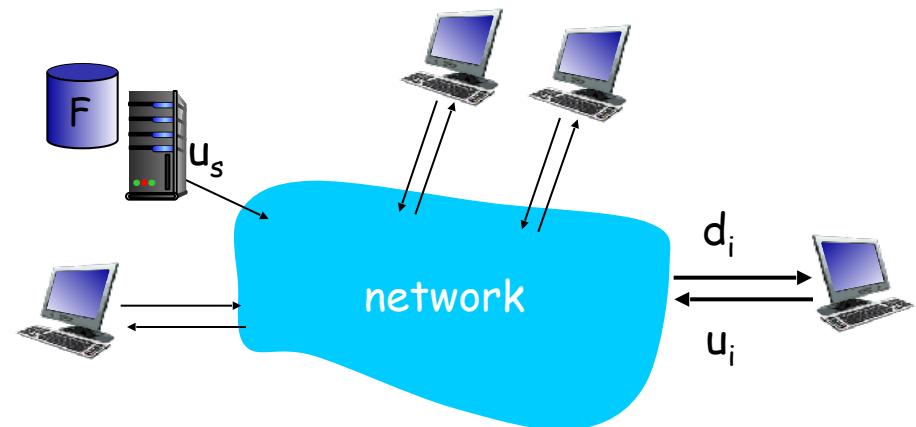
time to distribute F
to N clients using
client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in N

File distribution time: P2P

- **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- § **client:** each client must download file copy
 - max client download time: F/d_{\min}



- § **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$

time to distribute F

to N clients using $D_{P2P} > \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$

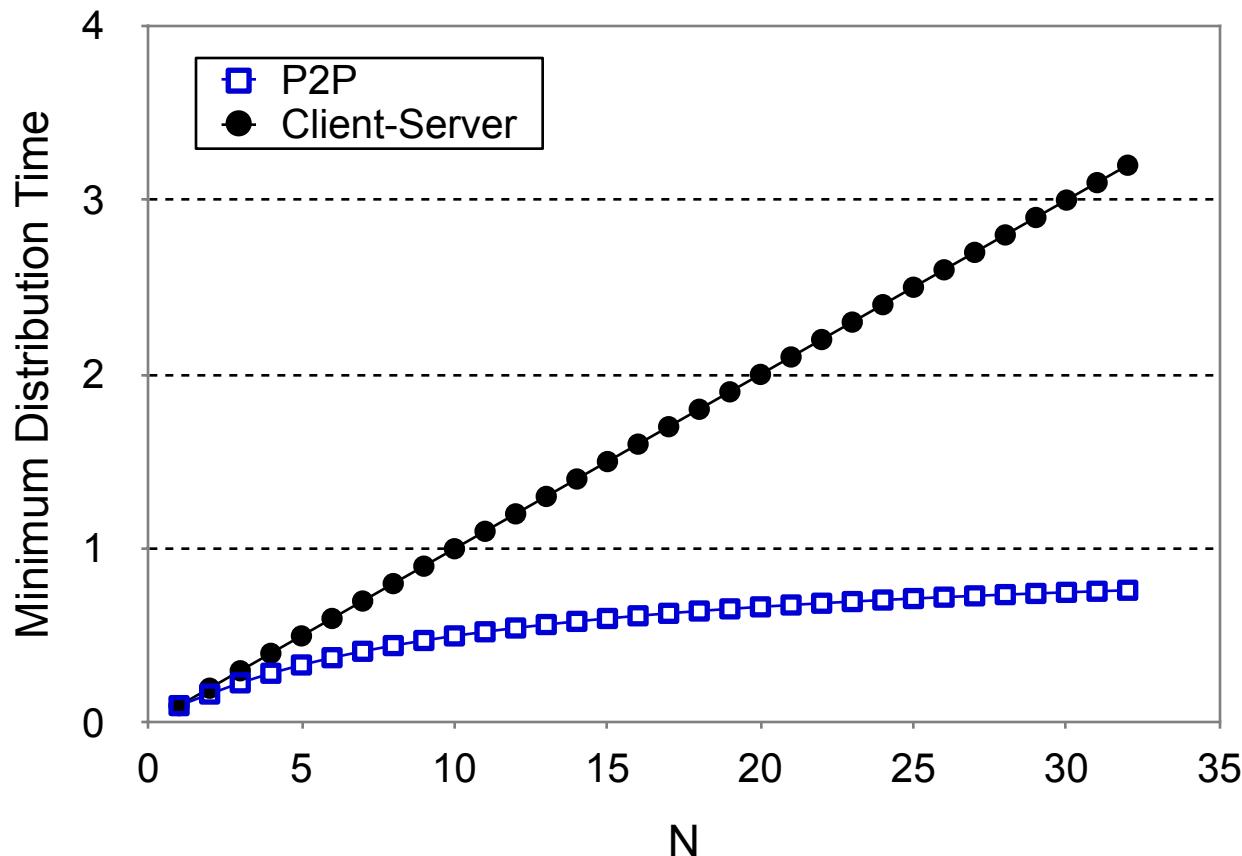
P2P approach

increases linearly in N ...

... but so does this, as each peer brings service capacity

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{\min} \geq u_s$



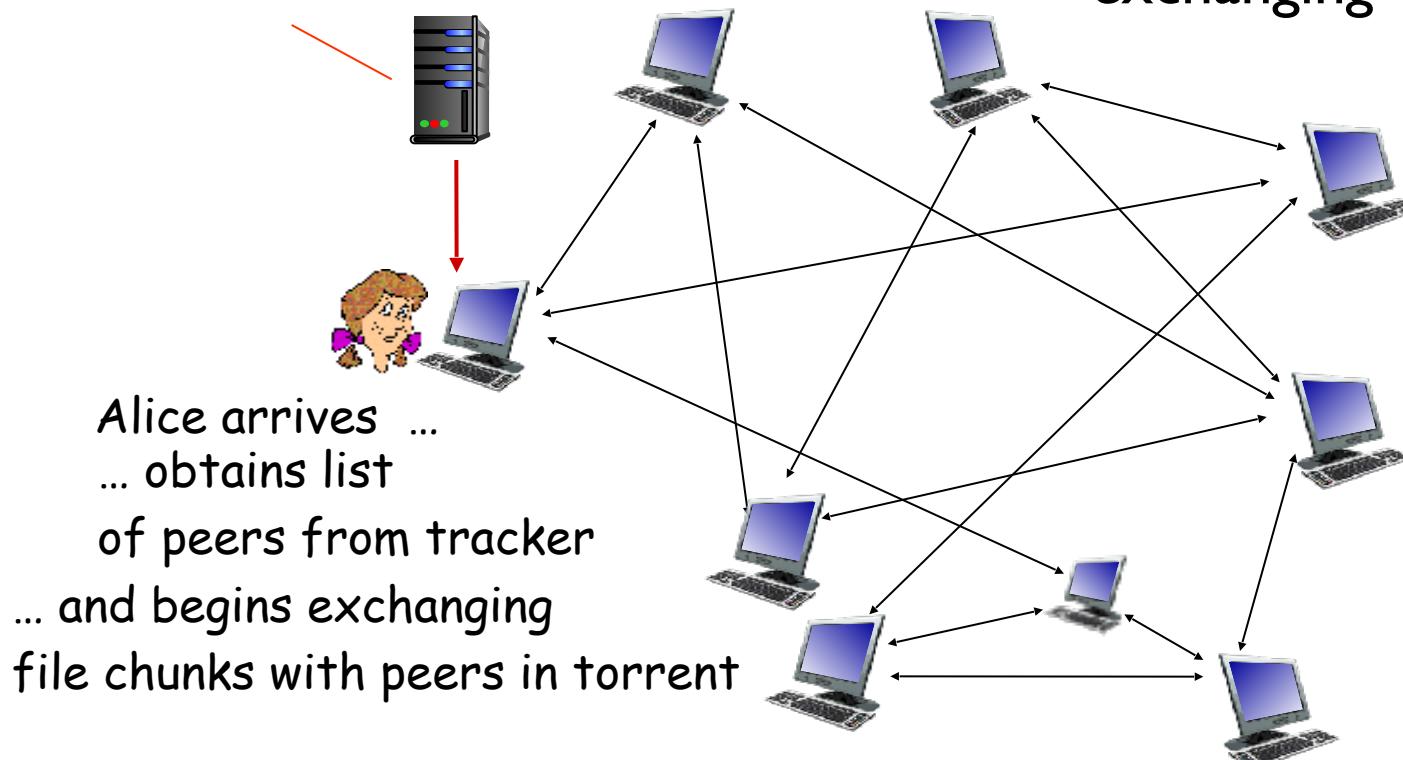
P2P file distribution: BitTorrent

§ file divided into chunks, e.g. of 256Kb each

§ peers in torrent send/receive file chunks

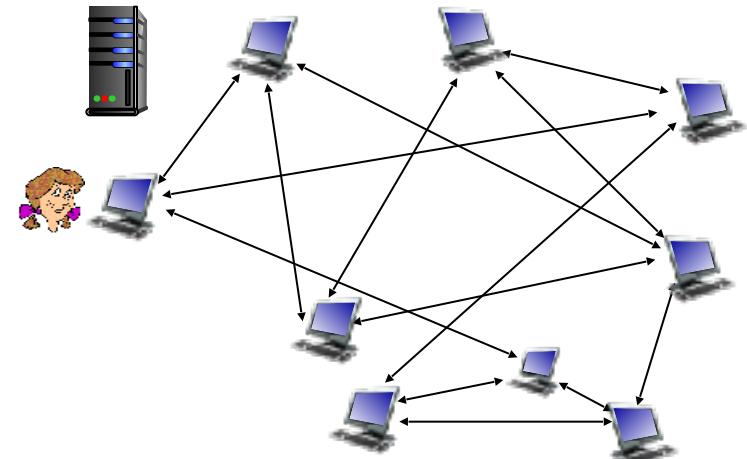
tracker: tracks peers
participating in torrent

torrent: group of peers
exchanging chunks of a file



P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)



- § while downloading, peer uploads chunks to other peers
- § peer may change peers with whom it exchanges chunks
- § **churn:** peers may come and go
- § once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

BitTorrent: requesting, sending file chunks

requesting chunks:

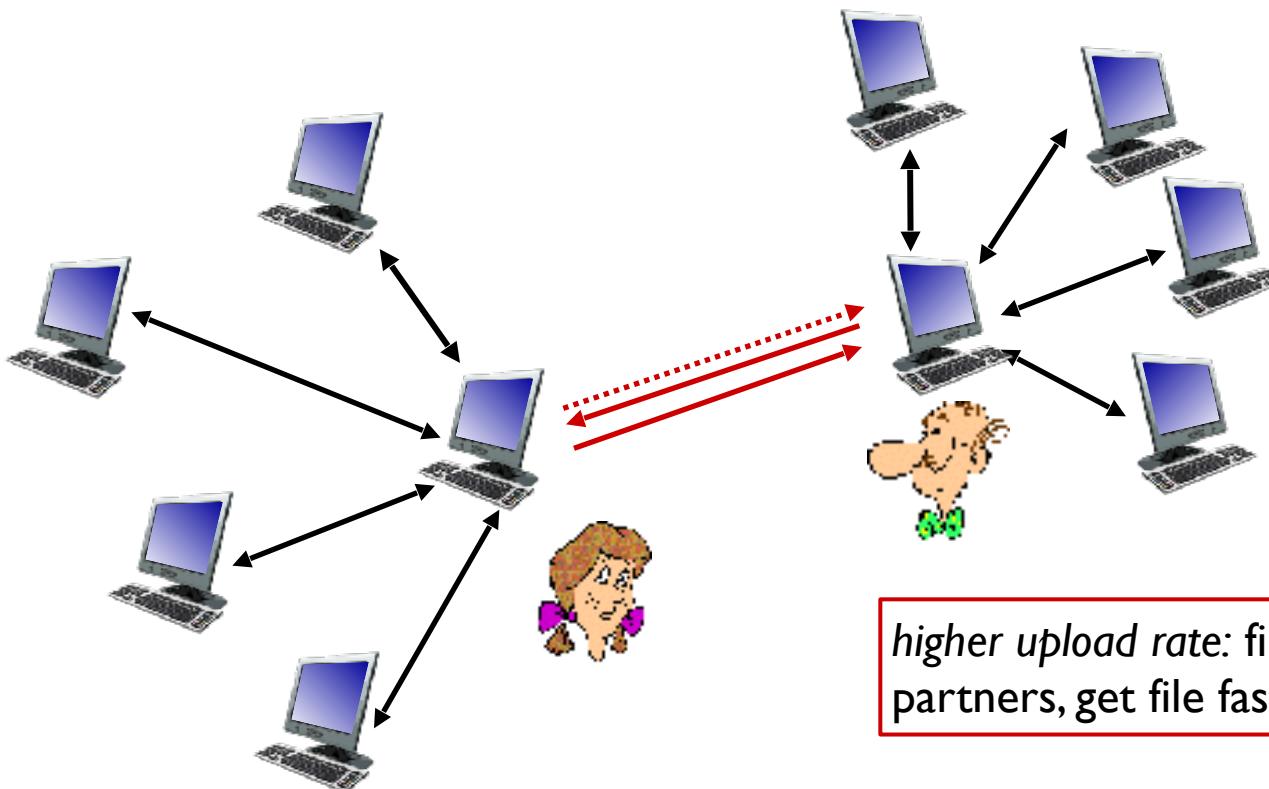
- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- § Alice sends chunks to those peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top peers every 10 secs
- § every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top peers

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top providers



higher upload rate: find better trading partners, get file faster !

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks (CDNs)

2.7 socket programming with UDP and TCP

Video Streaming and CDNs: context

§ video traffic: major consumer of Internet bandwidth

- Netflix, YouTube: 11%, 15% of global Internet traffic during pandemic stay-at-home orders
- ~2B YouTube users, ~193M Netflix users
- (last time slides were updated: 1B, 75M)

§ challenge: scale - how to reach ~2B users?

- single mega-video server won't work (why?)

§ challenge: heterogeneity

§ different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)

§ solution: distributed, application-level infrastructure



Multimedia: video

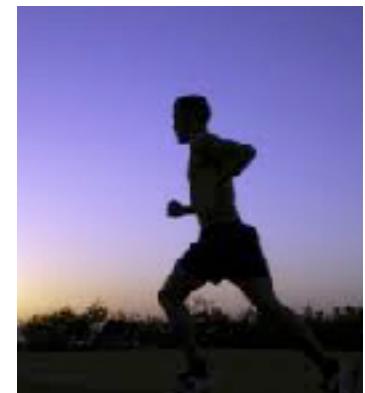
- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- digital image: array of pixels
 - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



frame $i+1$

Multimedia: video

- § **CBR: (constant bit rate):** video encoding rate fixed
- § **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- § **examples:**
 - MPEG 1 (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
(Columbia patent!
EE Prof Dimitris Anastassiou!!)
 - MPEG4 (often used in Internet, < 1 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame

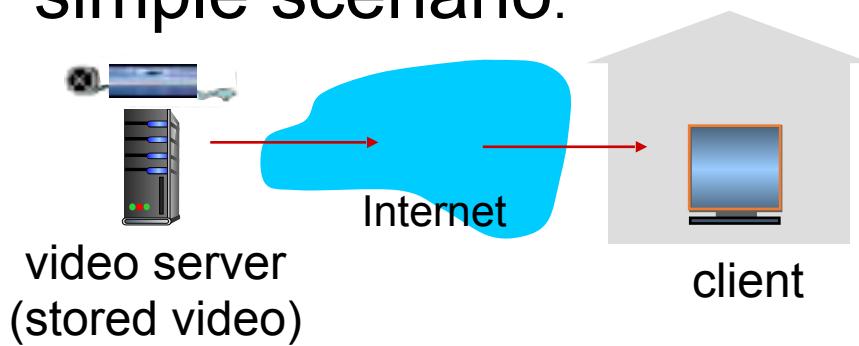
temporal coding example:
instead of sending complete frame at $i+1$, send only differences from frame i



frame

Streaming stored video

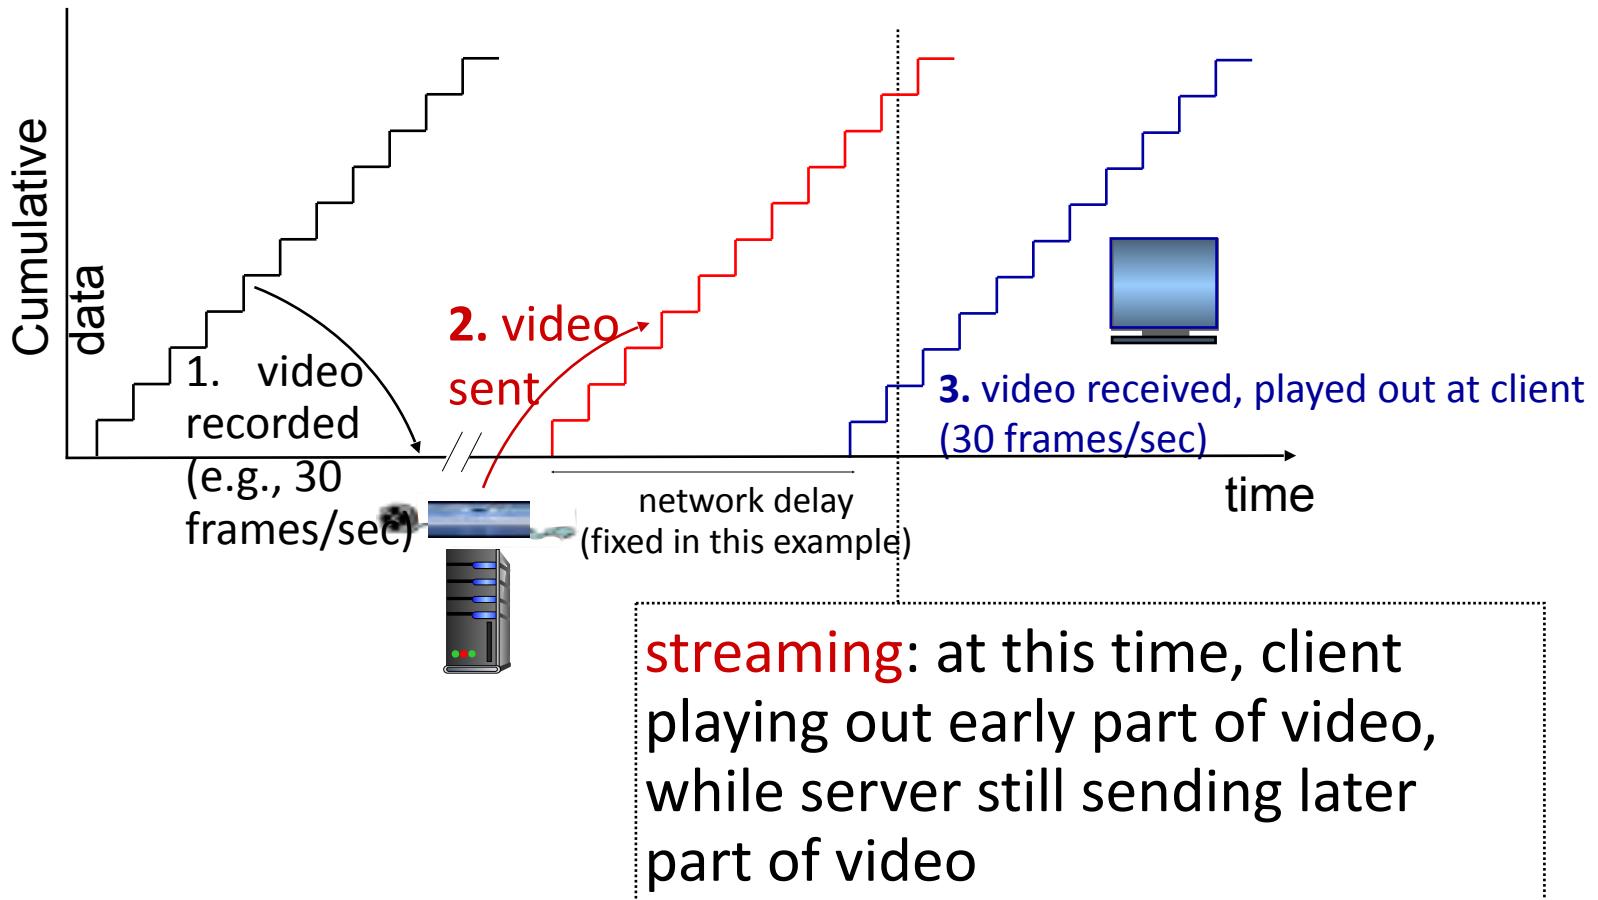
simple scenario:



Main challenges:

- server-to-client bandwidth will *vary* over time, with changing network congestion levels (in house, access network, network core, video server)
- packet loss, delay due to congestion will delay playout, or result in poor video quality

Streaming stored video

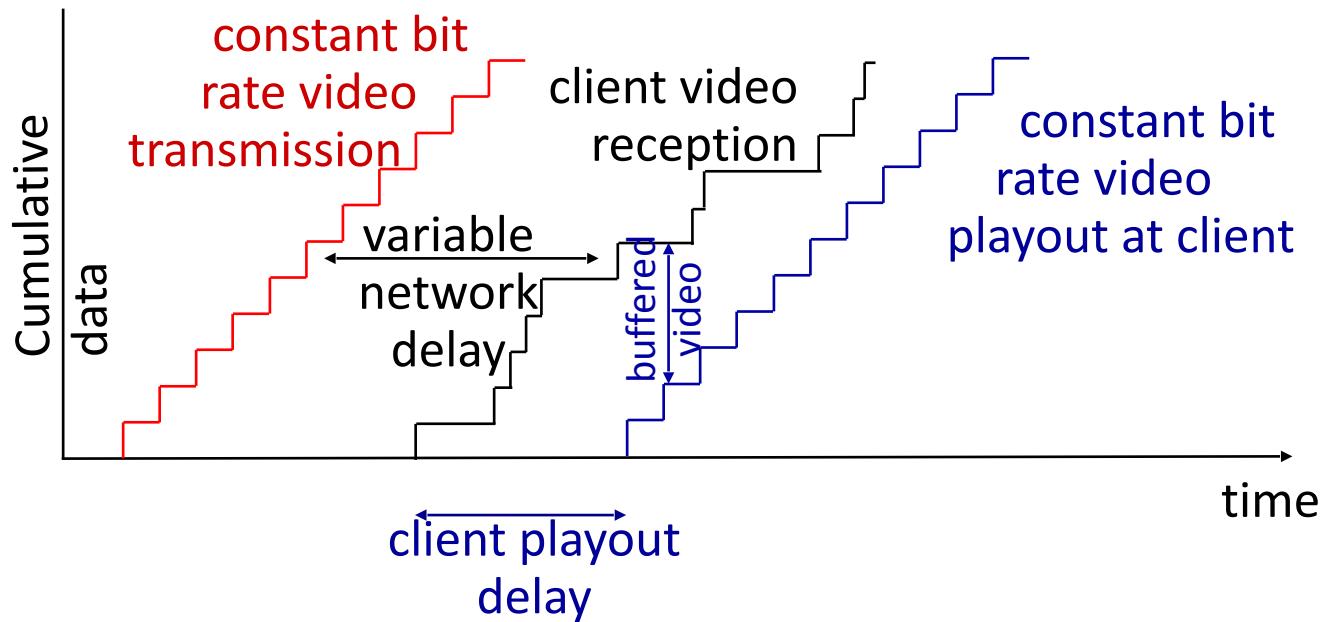


Streaming stored video: challenges

- **continuous playout constraint:** during client video playout, playout timing must match original timing
 - ... but **network delays are variable** (jitter), so will need **client-side buffer** to match continuous playout constraint
- other challenges:
 - client interactivity: pause, fast-forward, rewind, jump through video
 - video packets may be lost, retransmitted



Streaming stored video: playout buffering



- *client-side buffering and playout delay:* compensate for network-added delay, delay jitter

Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **server:**
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
 - *manifest file*: provides URLs for different chunks
 - why?
- **client:**
 - periodically measures server-to-client bandwidth (BW)
 - consulting manifest, requests one chunk at a time
 - chooses max encoding rate predicted to be sustainable given BW and/or playback buffer (current, possibly history)
 - can choose different coding rates at different points in time (depending on available bandwidth/buffer)

Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “*intelligence*” at client: client determines
 - *when* to request chunk (so that buffer starvation, or overflow does not occur)
 - and *when* to start playback
 - *what encoding rate* to request (higher quality when more bandwidth available)
- *(Some variants put intelligence at server or in cloud, tell client what to do)*

Day 12: CDNs and Transport Principles



CSEE 4119
Computer Networks
Ethan Katz-Bassett

 COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

Slides adapted from (and often identical to) slides from Kurose and Ross.

All material copyright 1996-2020

J.F Kurose and K.W. Ross, All Rights Reserved

Content distribution networks

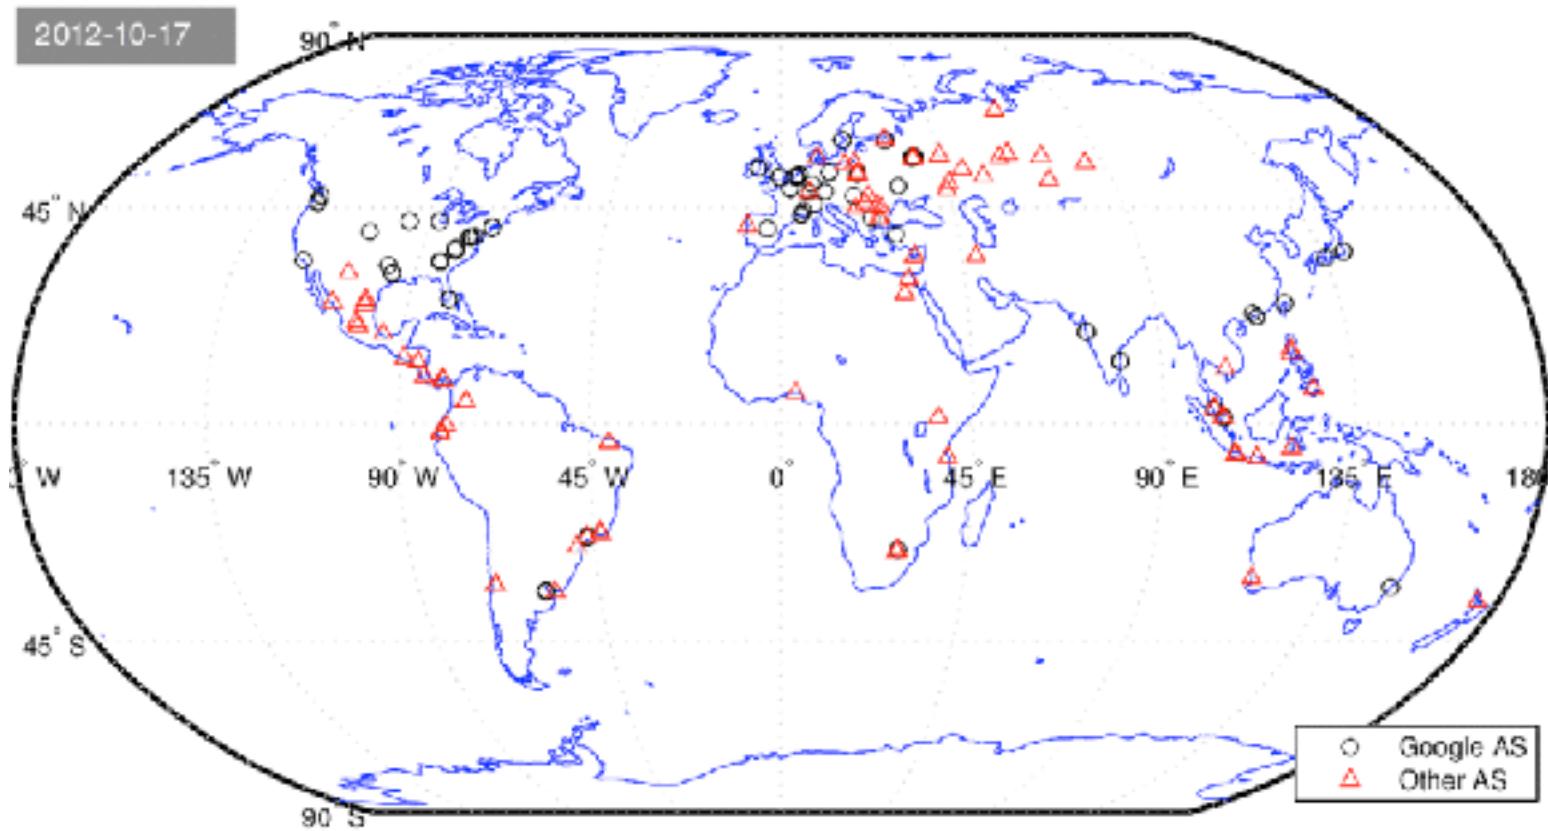
- ***challenge:*** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- ***option 1:*** single, large “mega-server”
 - single point of failure
 - point of network congestion
 - long path to distant clients
 - multiple copies of video sent over outgoing link

....quite simply: this solution *doesn't scale*

Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **option 2:** store/serve multiple copies of video chunks at multiple geographically distributed sites (**CDN**)
 - CDN steers client request to server that is “close” to client or has high available bandwidth)
 - **enter deep:** push CDN servers into many access networks
 - close to users
 - used by Akamai (3300 locations) and Google. Recently, Netflix and Facebook
 - **bring home:** smaller number (10’s) of larger clusters in POPs near (but not within) access networks
 - used by Limelight and most others

Google moved from “bring home” to “enter deep”



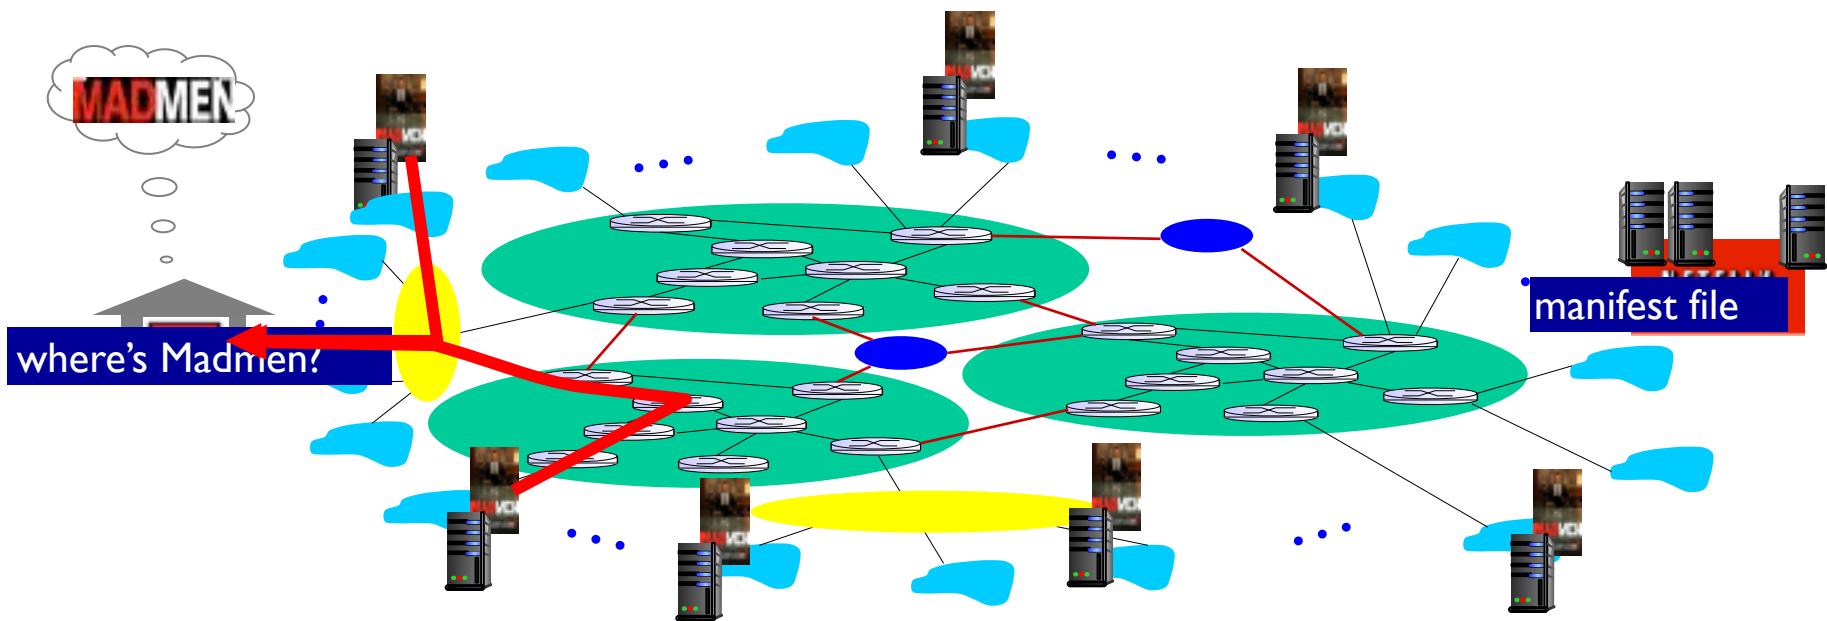
Our research exposed rapid growth in Google's deployment over a year

- Oct 2012: 200 sites in 60 countries
- Oct 2013: 1400 (7x) sites in 130+ (2.3x) countries

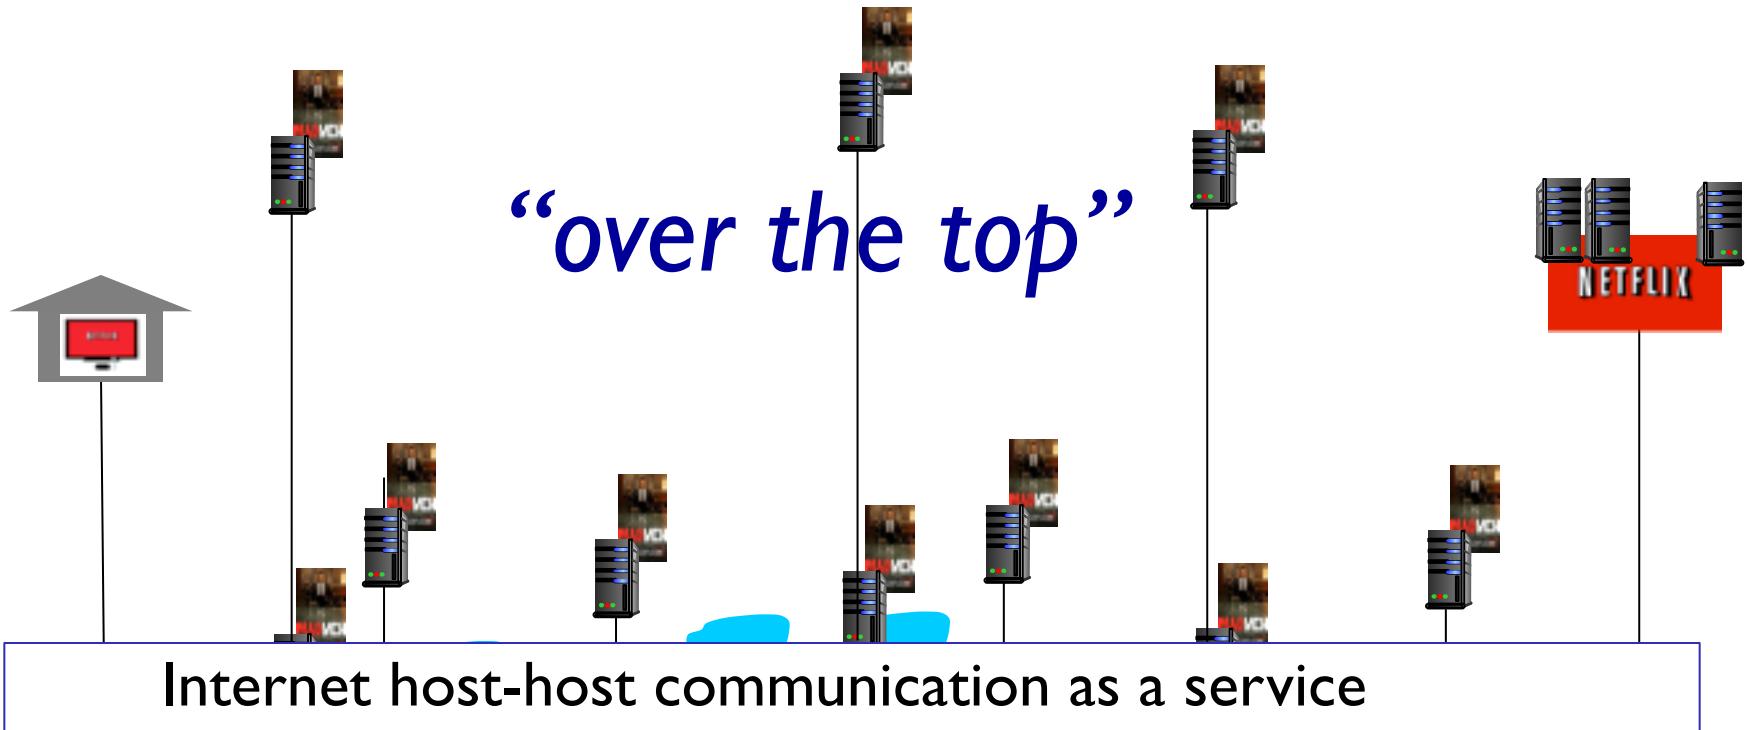
Content Distribution Networks (CDNs)

§ CDN: stores copies of content at CDN nodes

- e.g. Netflix stores copies of Mad Men
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Content Distribution Networks (CDNs)

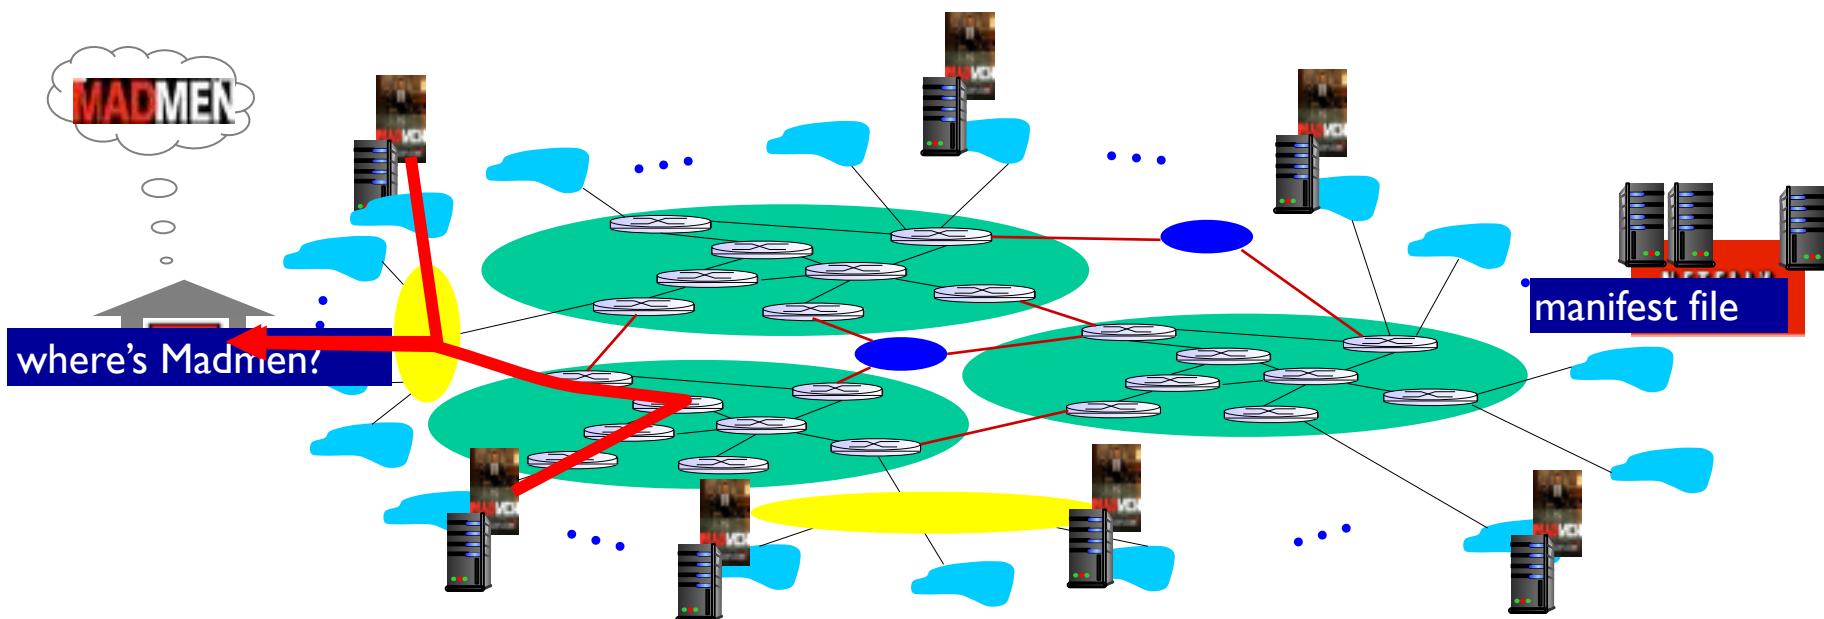


OTT challenges: coping with a congested Internet

- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

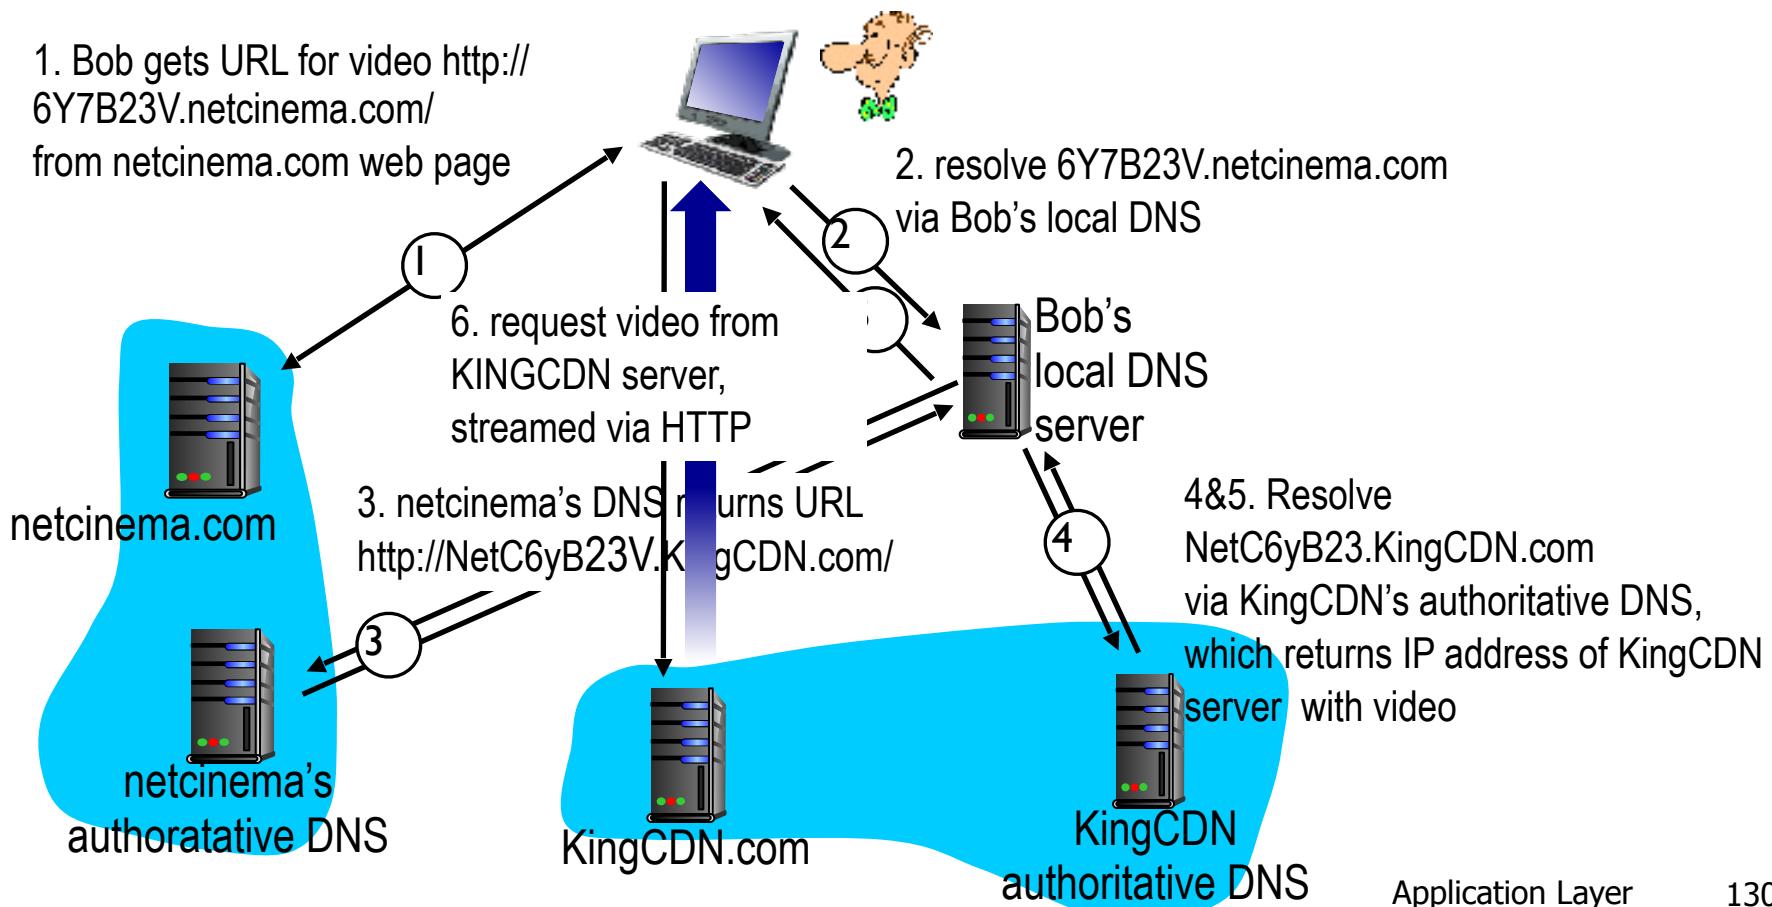
How to direct client to particular CDN server?

- 3 main approaches:
 - custom URL



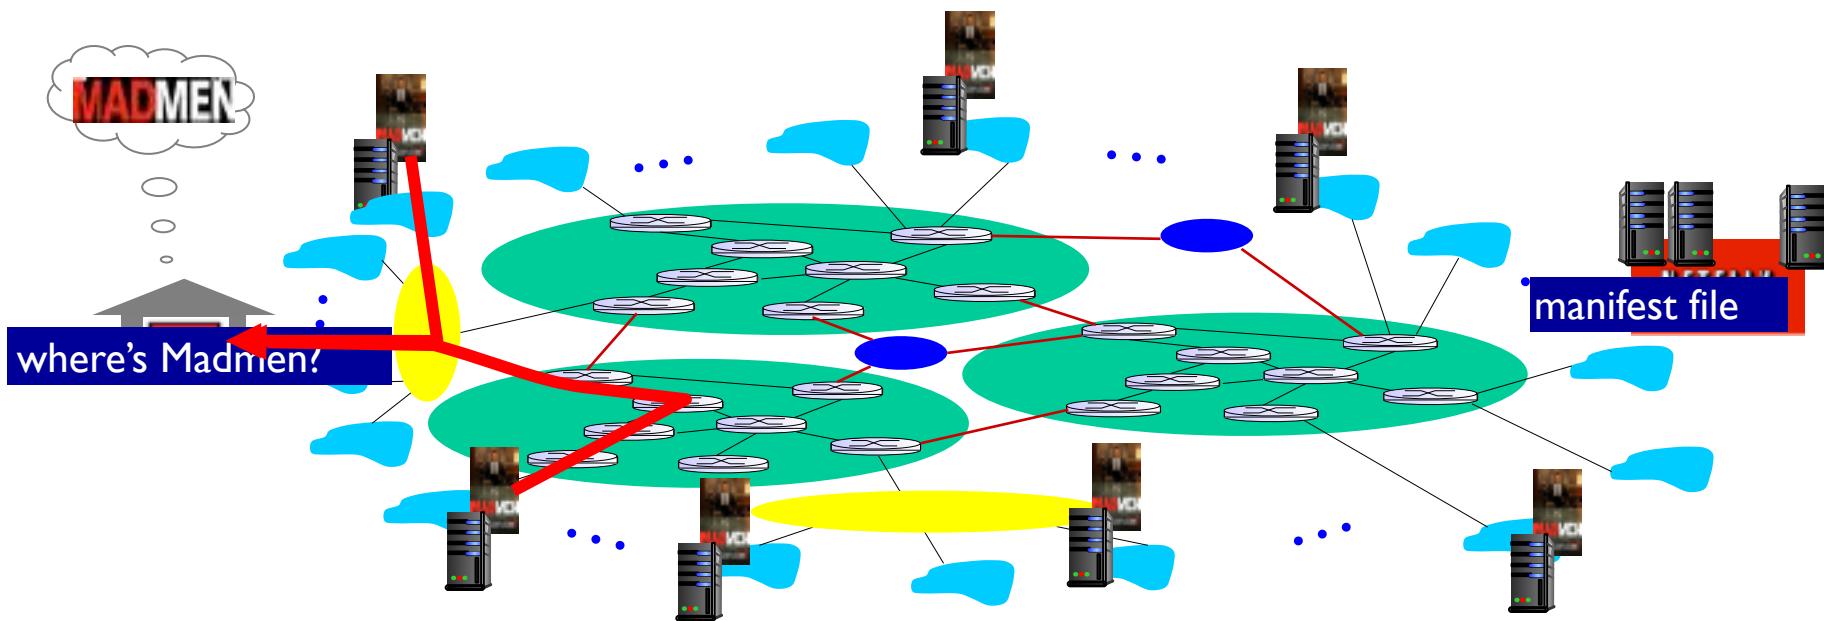
CDN content access: a closer look

Bob (client) requests video <http://6Y7B23V.netcinema.com/>
§video stored in CDN at <http://NetC6y&B23V.KingCDN.com/>



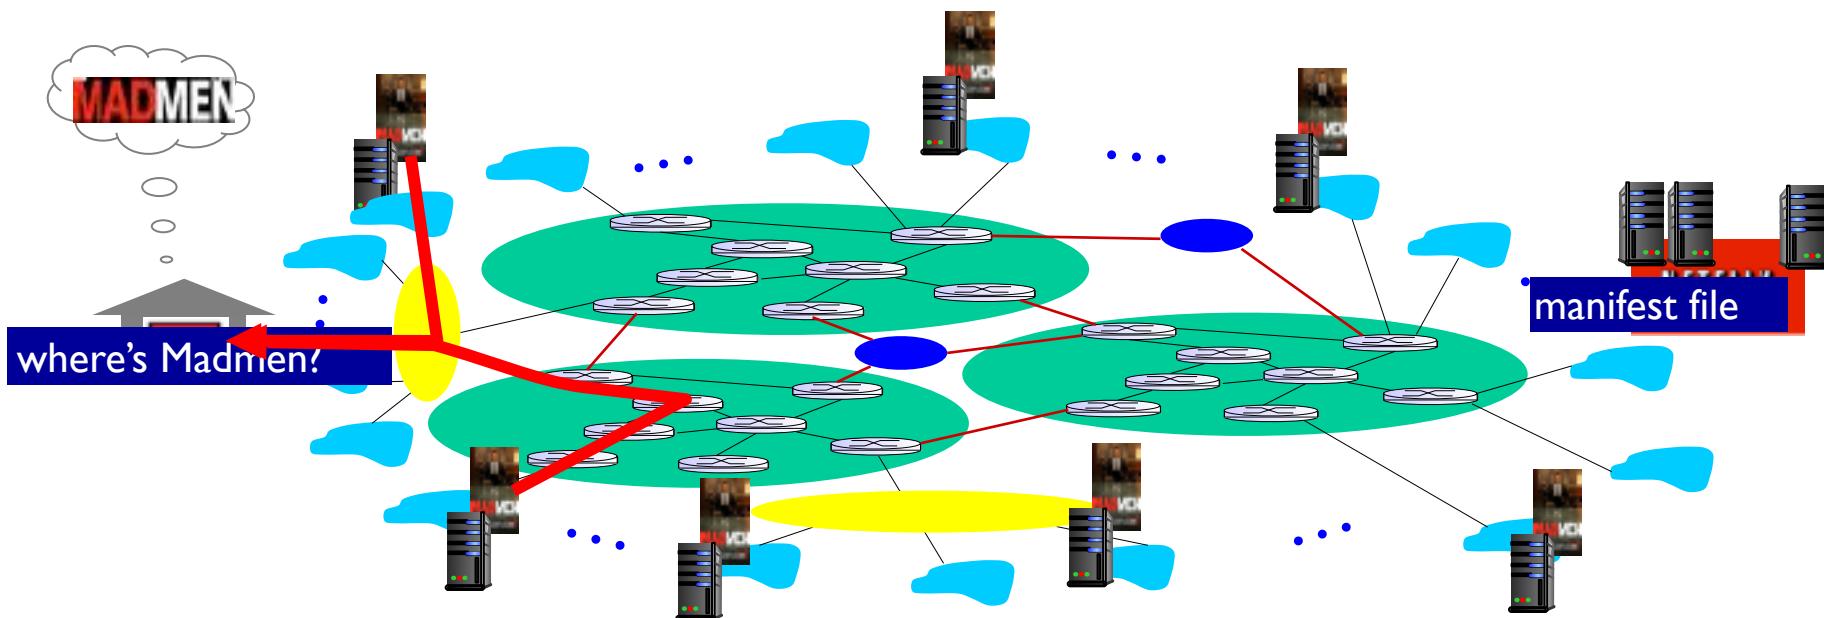
How to direct client to particular CDN server?

- 3 main approaches:
 - custom URLs
 - pros?
 - cons?



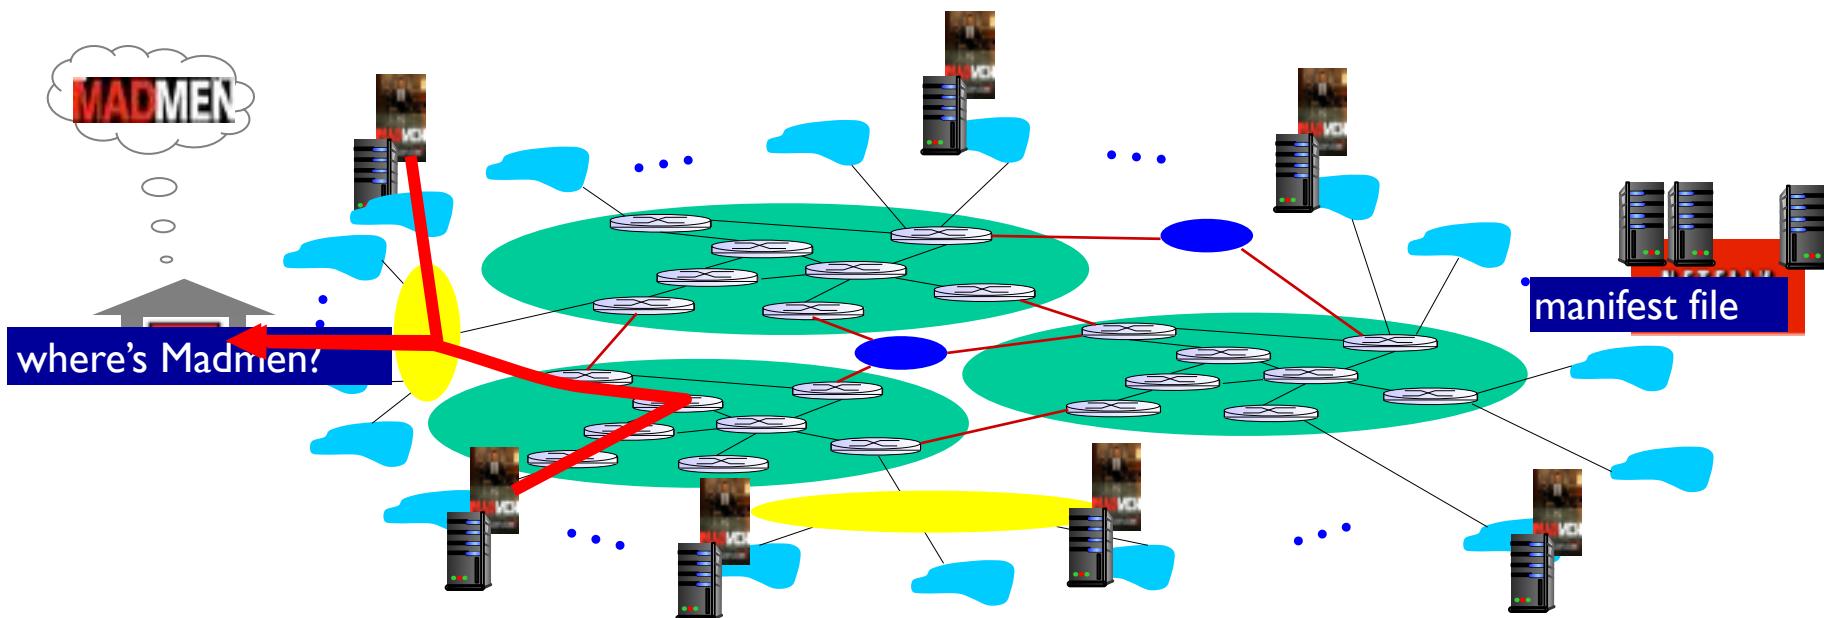
How to direct client to particular CDN server?

- 3 main approaches:
 - custom URLs
 - pros? precise per-client targeting
 - cons? not relevant on first page, since customization is embedded in URL in HTML

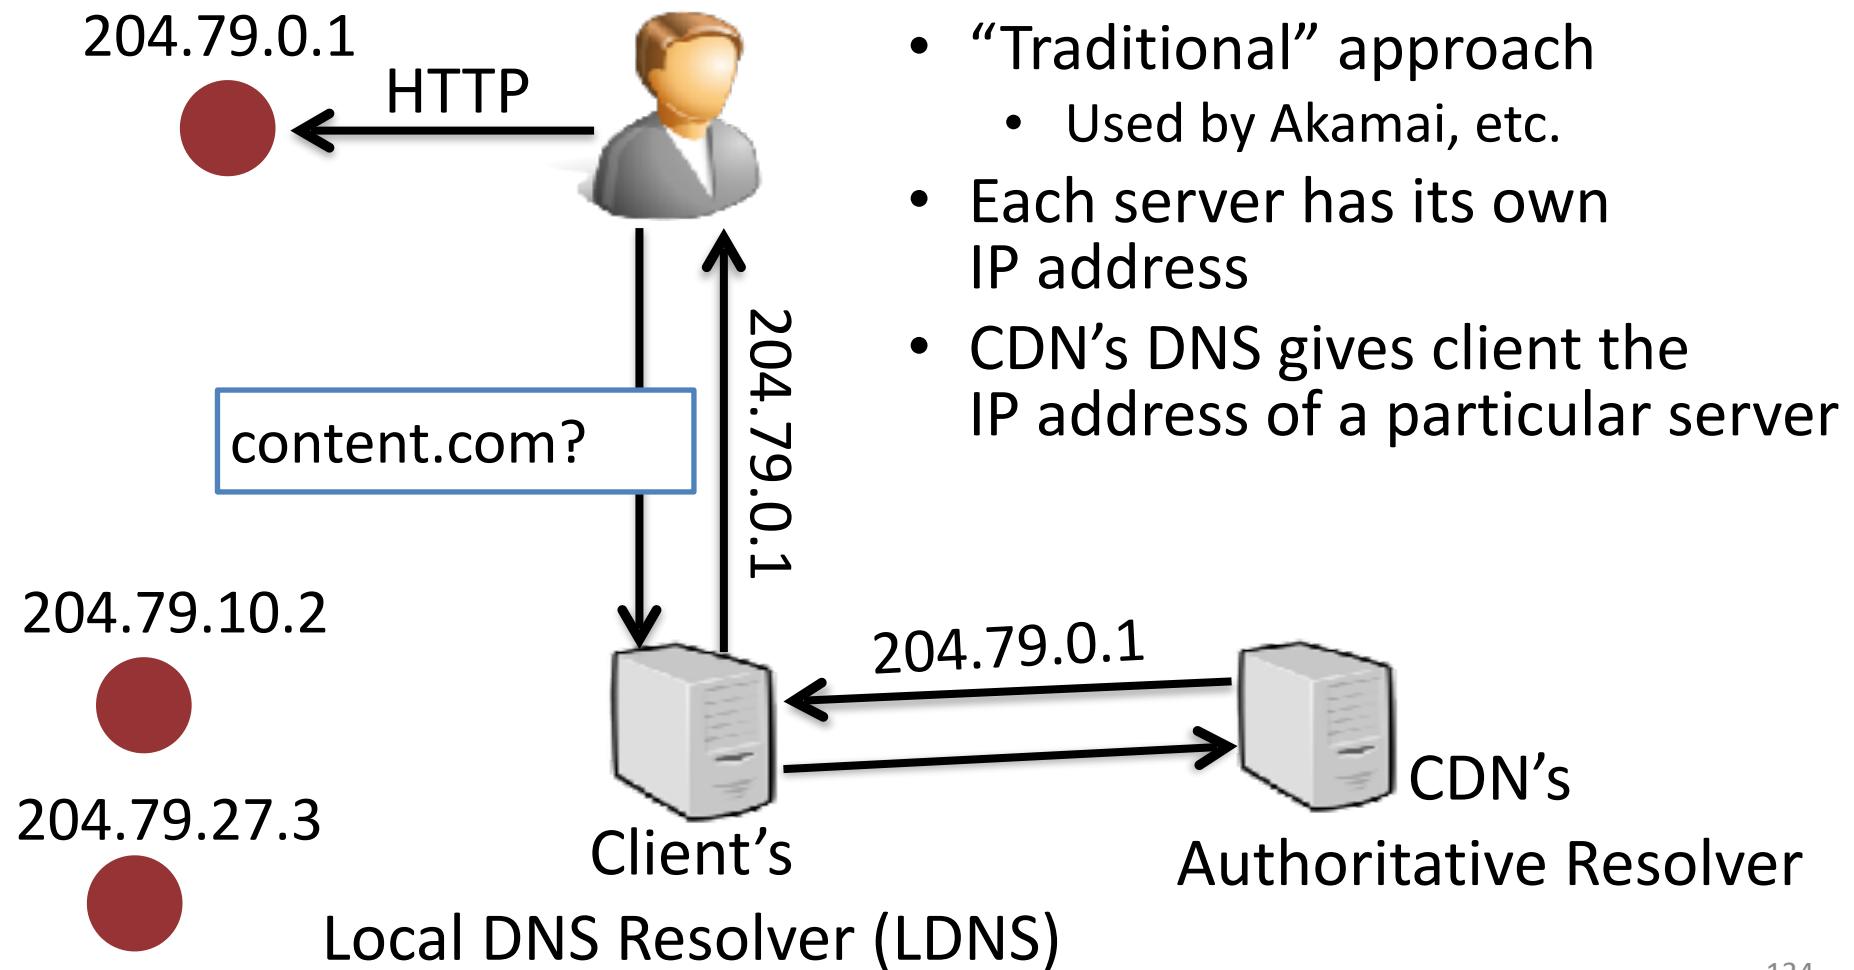


How to direct client to particular CDN server?

- 3 main approaches:
 - custom URLs
 - DNS redirection

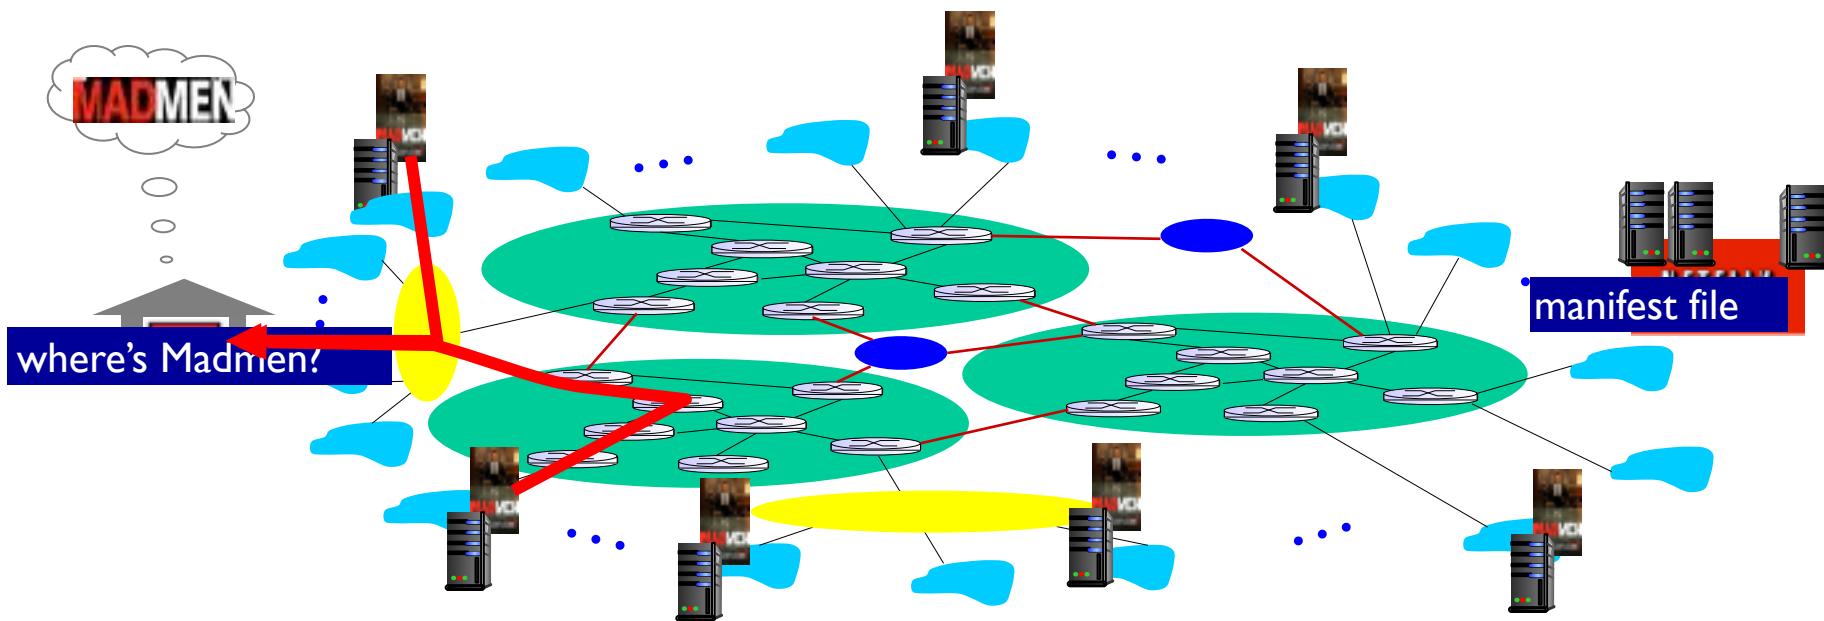


DNS Redirection CDN

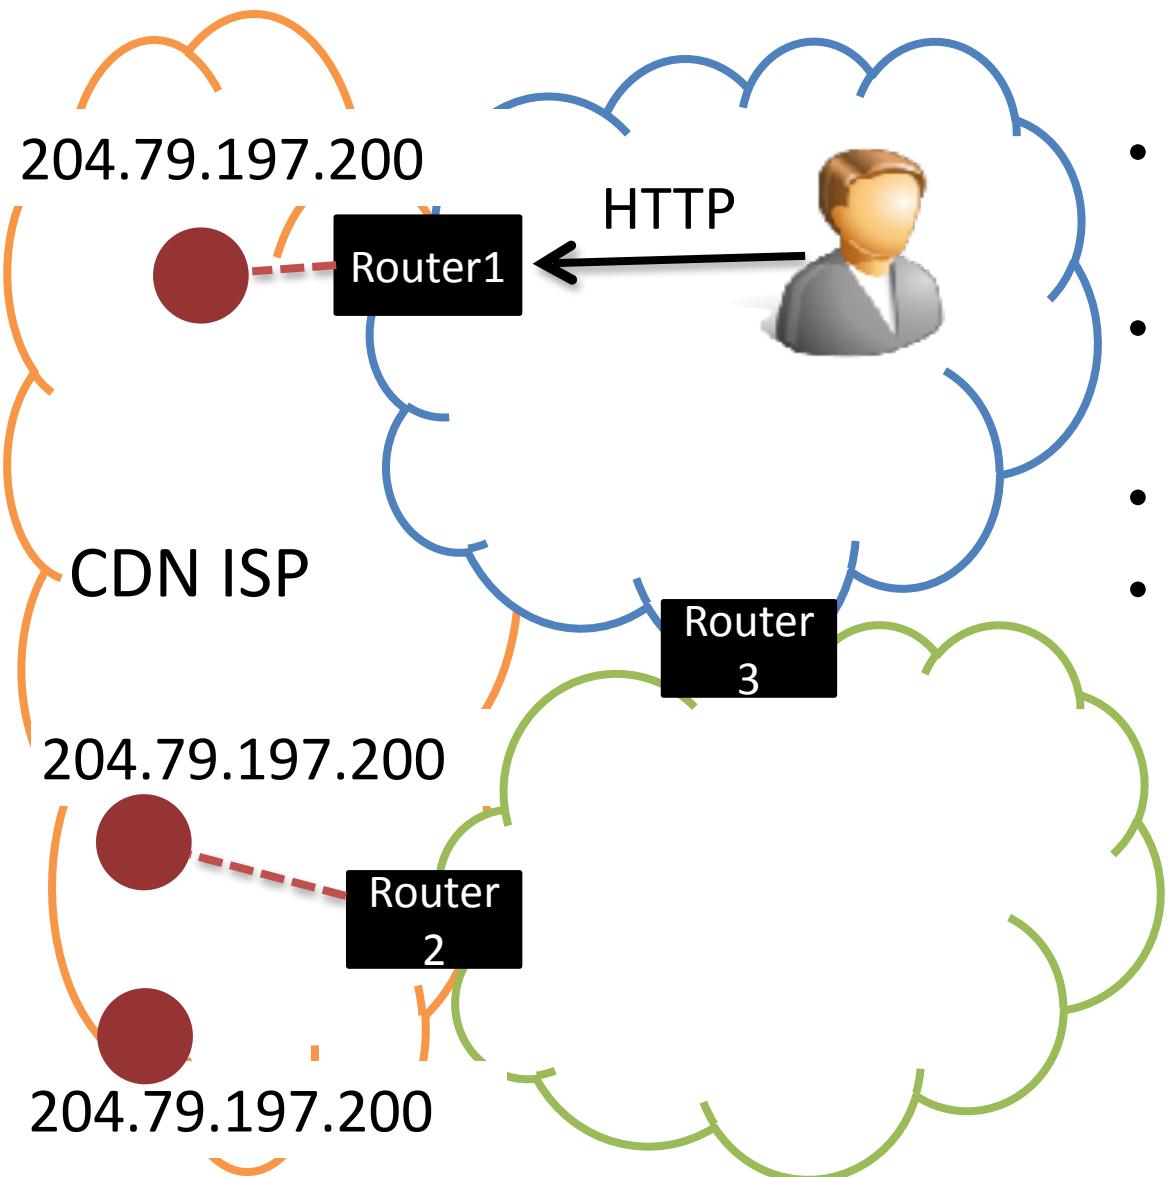


How to direct client to particular CDN server?

- 3 main approaches:
 - custom URLs
 - DNS redirection
 - anycast routing



Anycast CDN



- Gaining in popularity
 - EdgeCast, etc.
- All servers use same IP address
- BGP routing picks server
- *Run HTTP/TCP over anycast*

DNS

Anycast

Operational Complexity

- Complex: requires global traffic manager
- + Simple to deploy and operate
- + Naturally resilient against DDOS
- + High availability and fast fail over

Control

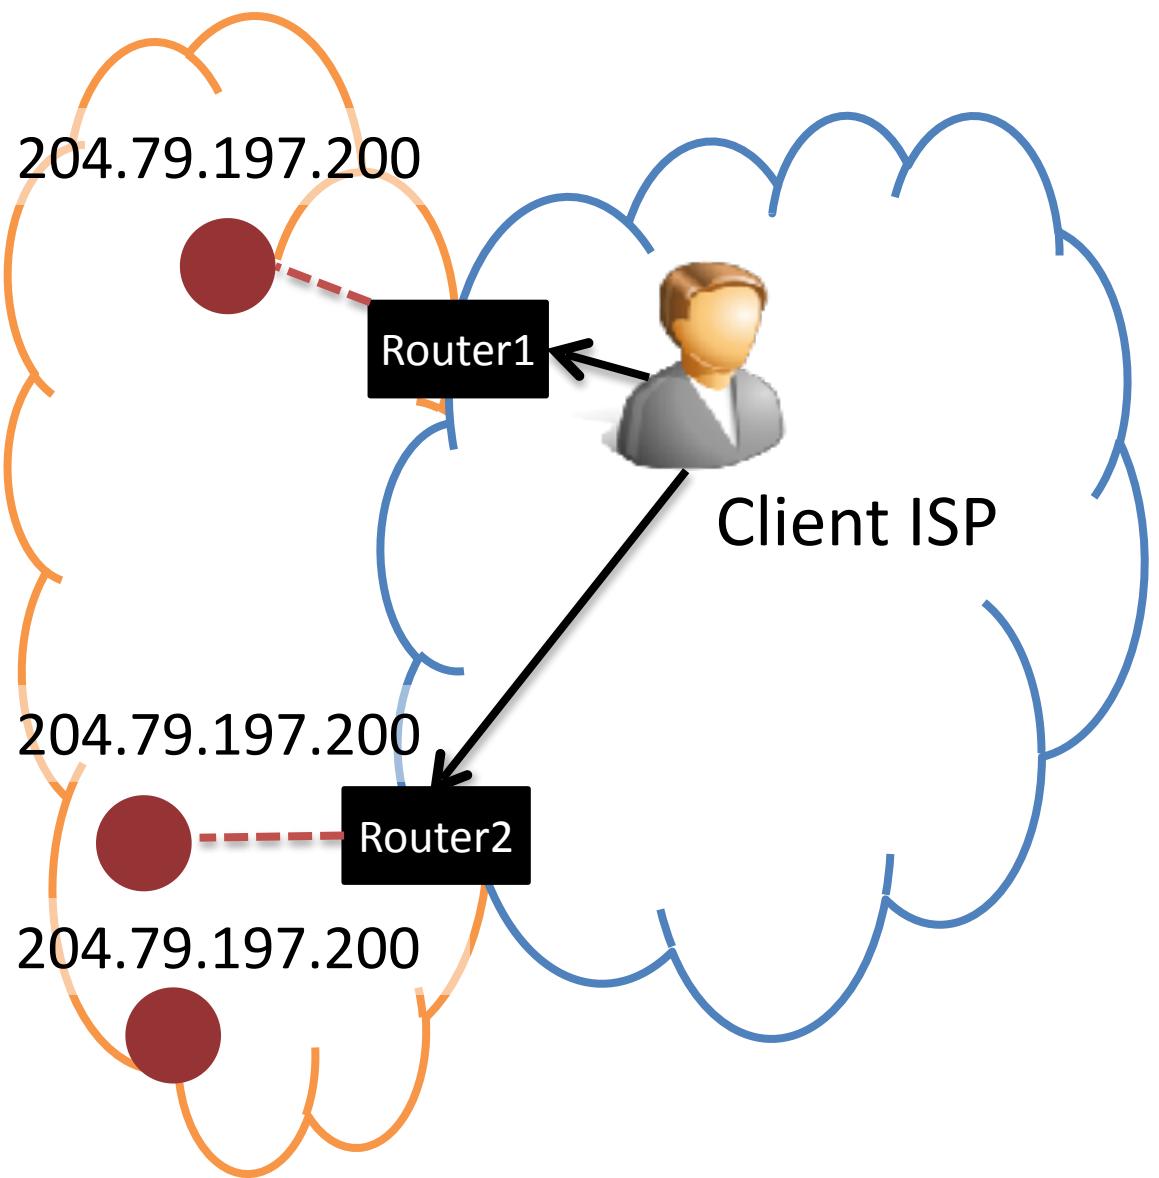
- + CDN-controlled
 - + Picks server of choice
 - + React quickly to:
 - Changes in load
 - Changes in performance

- controlled by distributed BGP routing

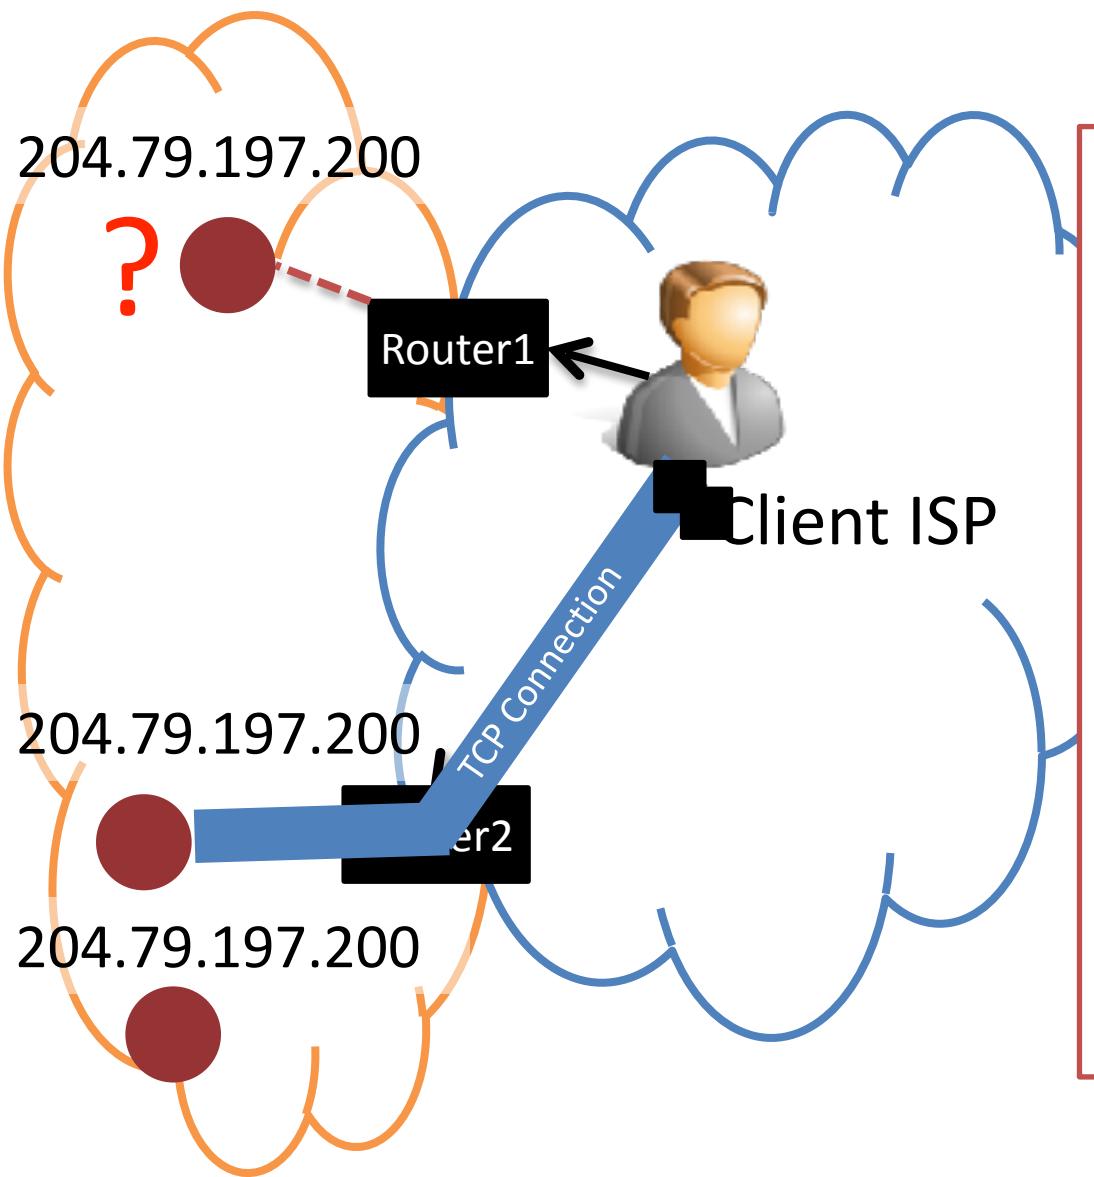
?

Granularity

Anycast pitfall: BGP may route to distant server



Anycast Pitfalls



- BGP may direct to distant server
- Server switches reset TCP connections

DNS

Operational Complexity

- Complex: requires global traffic manager

Anycast

Control

- + CDN-controlled
 - + Picks server of choice
 - + React quickly to:
 - Changes in performance
 - Changes in load

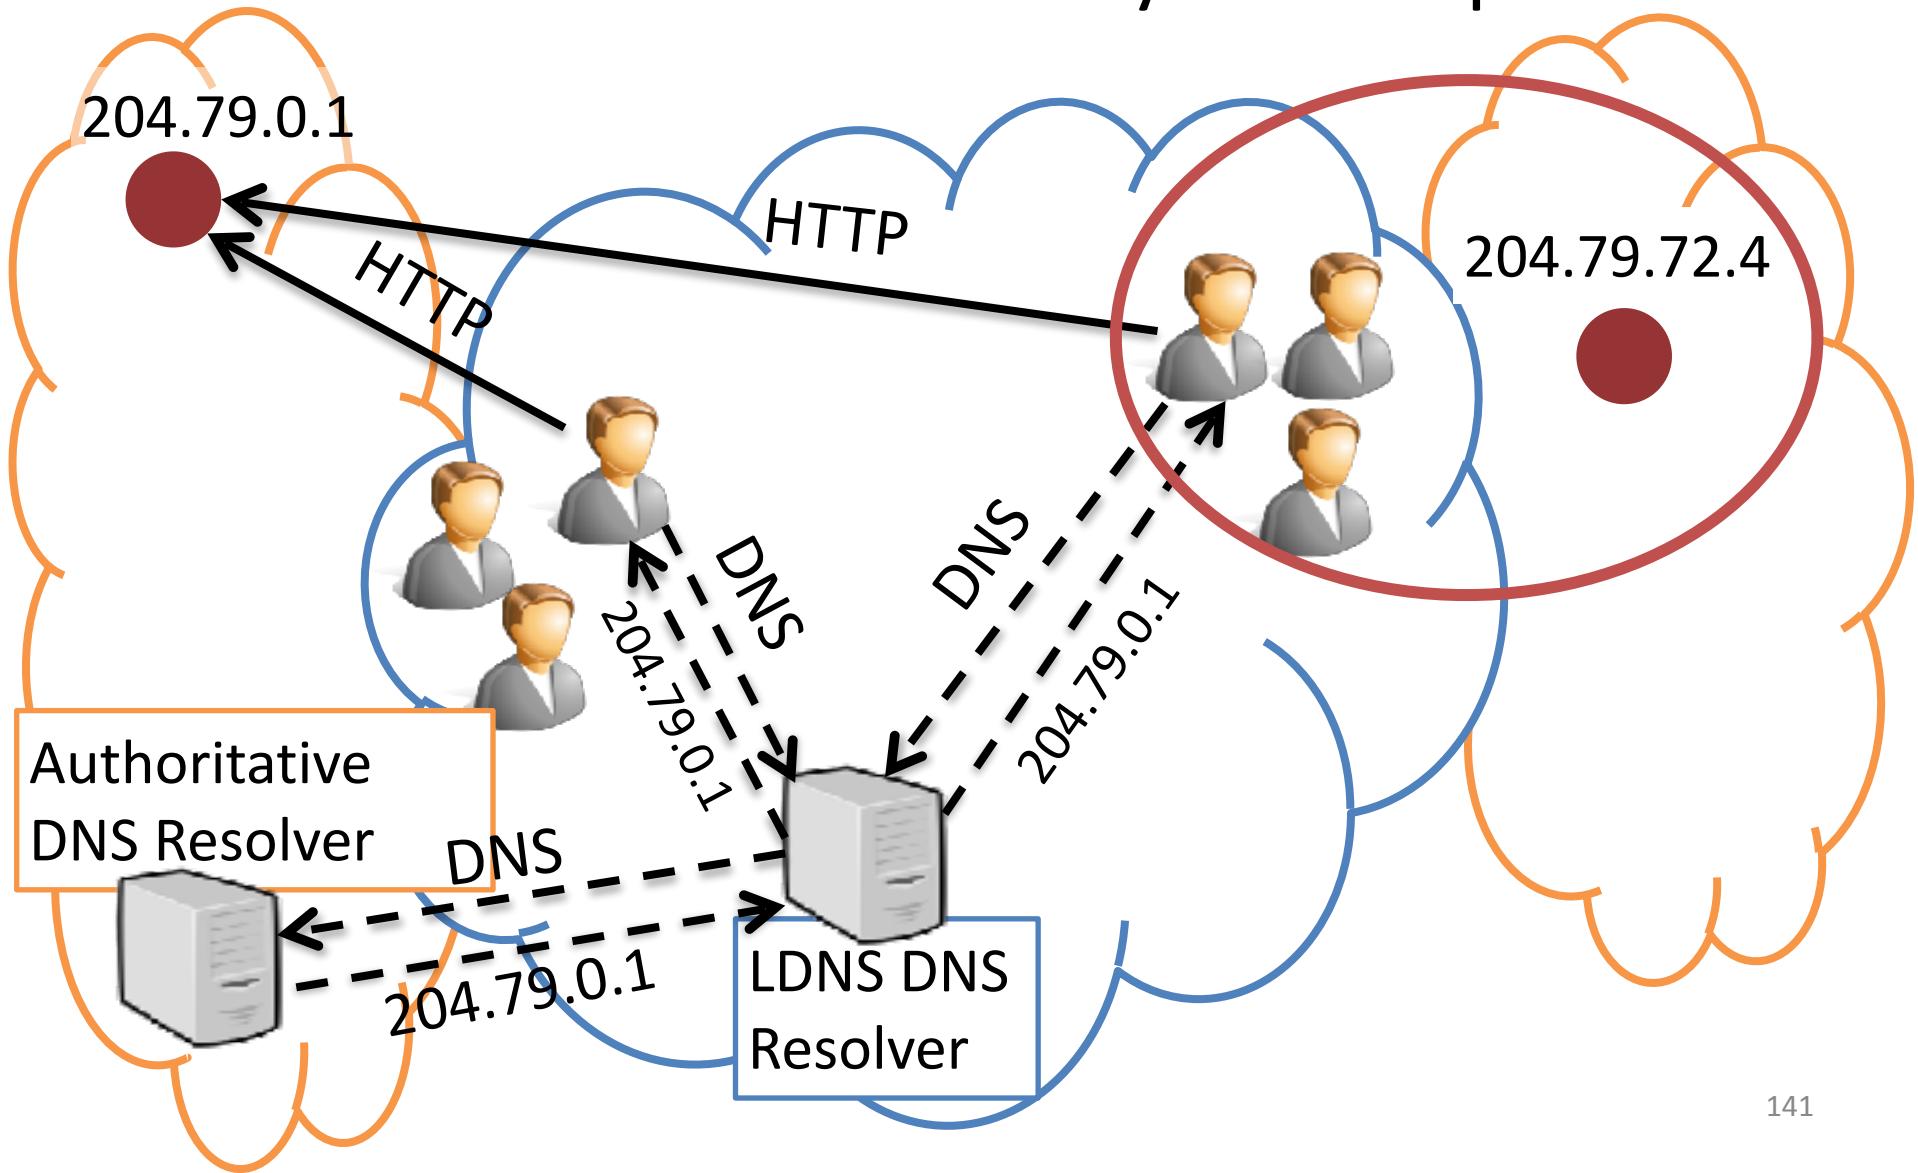
- BGP-controlled
 - Subject to vagaries of BGP routing
 - May not pick best
 - Hard to control load
 - TCP can break

Granularity

?

- + Per packet

DNS pitfall: Decisions are per LDNS, but clients of LDNS may be far apart



DNS

Anycast

Operational Complexity

- Complex: requires global traffic manager
- + Simple to deploy and operate
- + Naturally resilient against DDOS
- + High availability and fast fail over

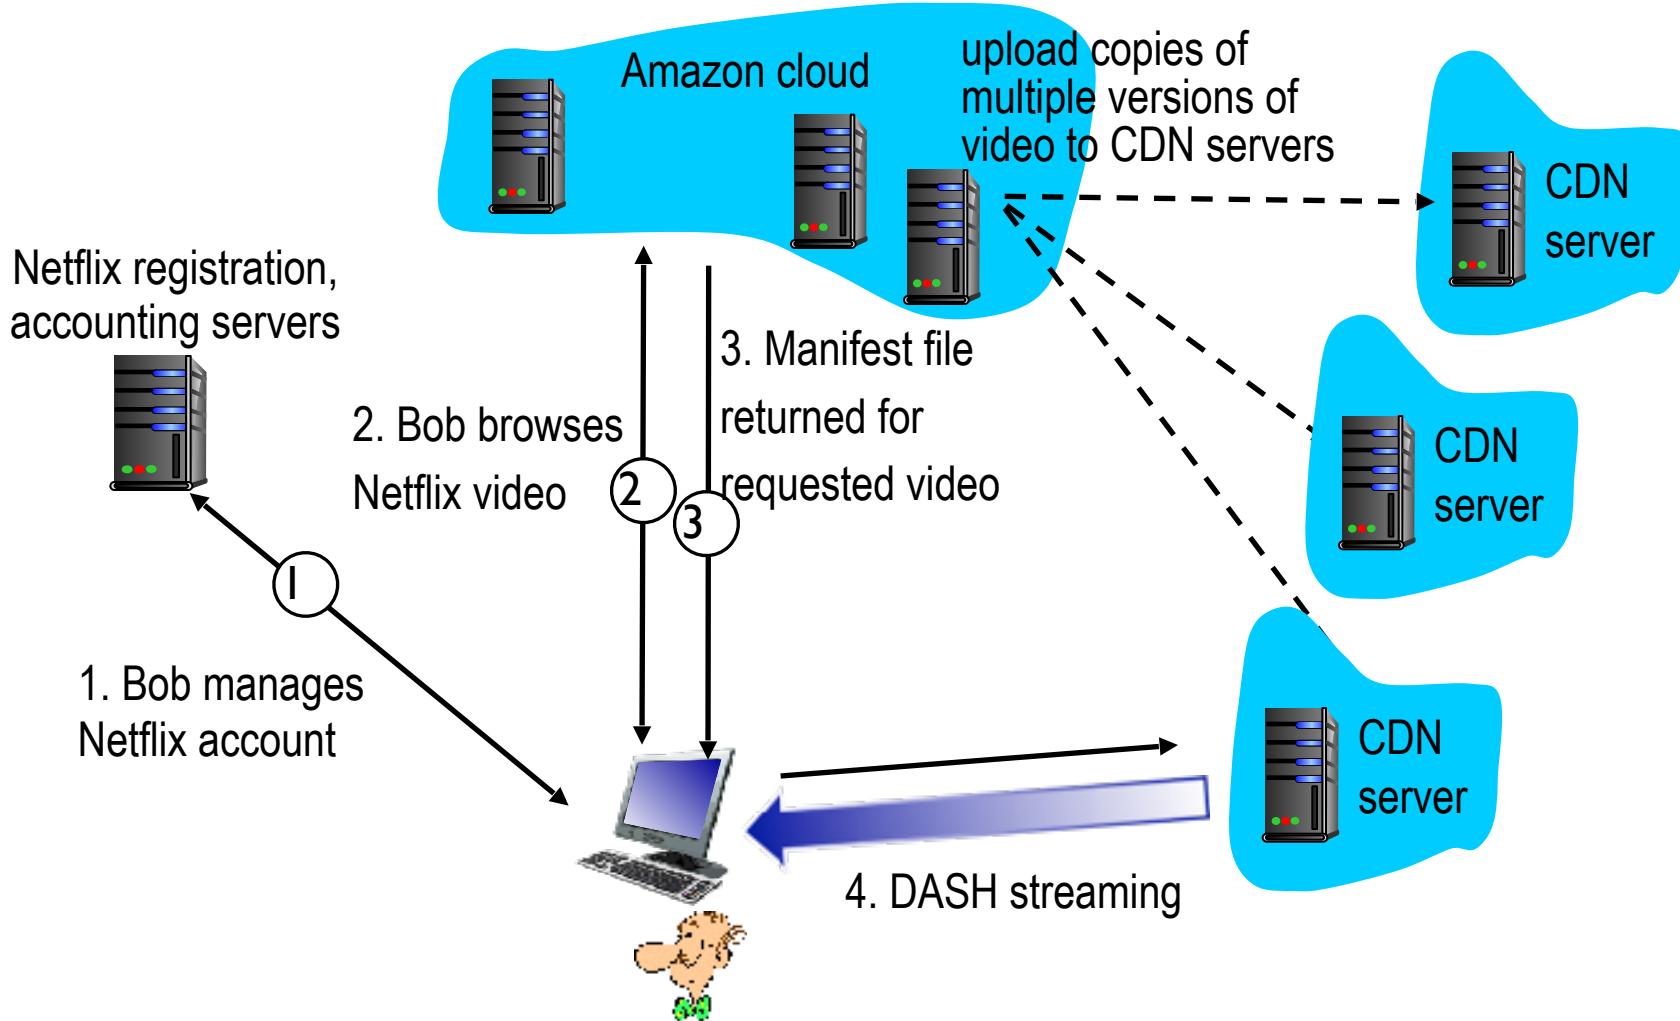
Control

- + CDN-controlled
 - + Picks server of choice
 - + React quickly to:
 - Changes in performance
 - Changes in load
- BGP-controlled
 - Subject to vagaries of BGP routing
 - May not pick nearest server
 - Hard to control load

Granularity

- Per LDNS can be poor representative of clients
- + Per packet

Case study: Netflix



Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Chapter 2: summary

our study of network apps now complete!

- application architectures
 - client-server
 - P2P
 - application service requirements:
 - reliability, bandwidth, delay
 - Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- § specific protocols:
- HTTP
 - SMTP, POP, IMAP
 - DNS
 - P2P: BitTorrent
- § video streaming, CDNs
- § socket programming:
TCP, UDP sockets

Chapter 2: summary

most importantly: learned about protocols!

- typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- message formats:
 - *headers*: fields giving info about data
 - *data*: info(payload) being communicated

important themes:

- § control vs. messages
 - in-band, out-of-band
- § centralized vs. decentralized
- § stateless vs. stateful
- § reliable vs. unreliable message transfer
- § “complexity at network edge”



**DO NOT SHARE
SLIDES AND CLASS MATERIALS
ON ONLINE SITES**

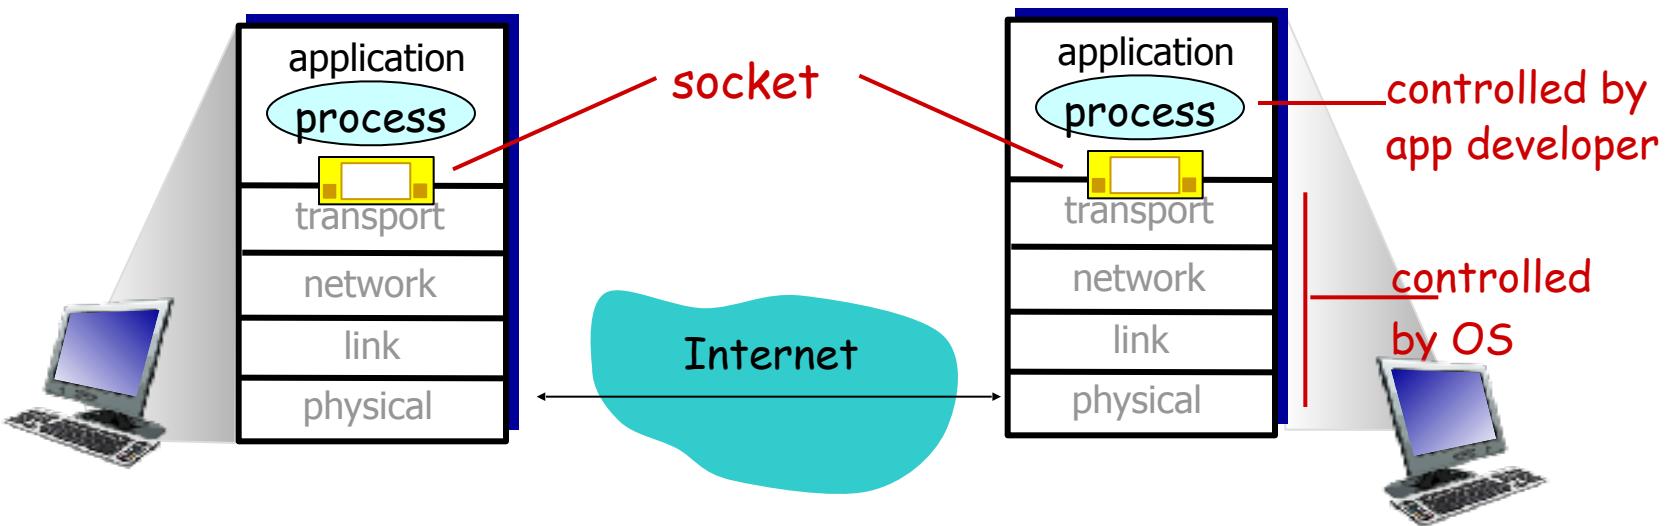
Course Hero

Uploading course materials to sites such as CourseHero, Chegg or Github is academic misconduct at Columbia (see [pg 10](#) of [Columbia guide](#)).

Socket programming

goal: learn how to build client/server applications that communicate using sockets

socket: door between application process and end-end-transport protocol



Socket programming

Two socket types for two transport services:

- **UDP**: unreliable datagram
- **TCP**: reliable, byte stream-oriented

Application Example:

1. client reads a line of characters (data) from its keyboard and sends data to server
2. server receives the data and converts characters to uppercase
3. server sends modified data to client
4. client receives modified data and displays line on its screen

Socket programming with UDP

UDP: no “connection” between client & server

- no handshaking before sending data
- sender explicitly attaches IP destination address and port # to each packet
- receiver extracts sender IP address and port# from received packet

UDP: transmitted data may be lost or received out-of-order

Application viewpoint:

- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

Client/server socket interaction: UDP

server (running on serverIP)

```
create socket, port= x:  
serverSocket =  
socket(AF_INET,SOCK_DGRAM)
```

read datagram from
serverSocket

write reply to
serverSocket
specifying
client address,
port number

client

```
create socket:  
clientSocket =  
socket(AF_INET,SOCK_DGRAM)
```

Create datagram with server IP and
port=x; send datagram via
clientSocket

read datagram from
clientSocket
close
clientSocket

Example app: UDP client

Python UDPClient

```
include Python's socket  
library → from socket import *  
  
create UDP socket for  
server → clientSocket = socket(AF_INET,  
                                SOCK_DGRAM)  
  
get user keyboard  
input → message = raw_input('Input lowercase sentence:')  
  
Attach server name, port → clientSocket.sendto(message.encode(),  
to message; send into  
socket → (serverName, serverPort))  
  
read reply characters → modifiedMessage, serverAddress =  
from  
socket into string → clientSocket.recvfrom(2048)  
  
print out received string → print modifiedMessage.decode()  
and close socket → clientSocket.close()
```

Example app: UDP server

Python UDPServer

```
from socket import *
serverPort = 12000
create UDP socket ----->serverSocket = socket(AF_INET, SOCK_DGRAM)
bind socket to local port number 12000 ----->serverSocket.bind(('', serverPort))
                                         print ("The server is ready to receive")
loop forever ----->while True:
Read from UDP socket into message, getting client's address (client IP and port) -----> message, clientAddress = serverSocket.recvfrom(2048)
                                         modifiedMessage = message.decode().upper()
send upper case string back to this client ----->serverSocket.sendto(modifiedMessage.encode(),
                                         clientAddress)
```

Socket programming with TCP

client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- *when client creates socket:* client TCP establishes connection to server TCP

- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
 - allows server to talk with multiple clients
 - source port numbers used to distinguish clients (more in Chap 3)

application viewpoint:

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

Client/server socket interaction: TCP

server (running on hostid)

create socket,
port=x, for incoming
request:
`serverSocket = socket()`

wait for incoming
connection request
`connectionSocket = serverSocket.accept()`

read request from
`connectionSocket`

write reply to
`connectionSocket`
close
`connectionSocket`

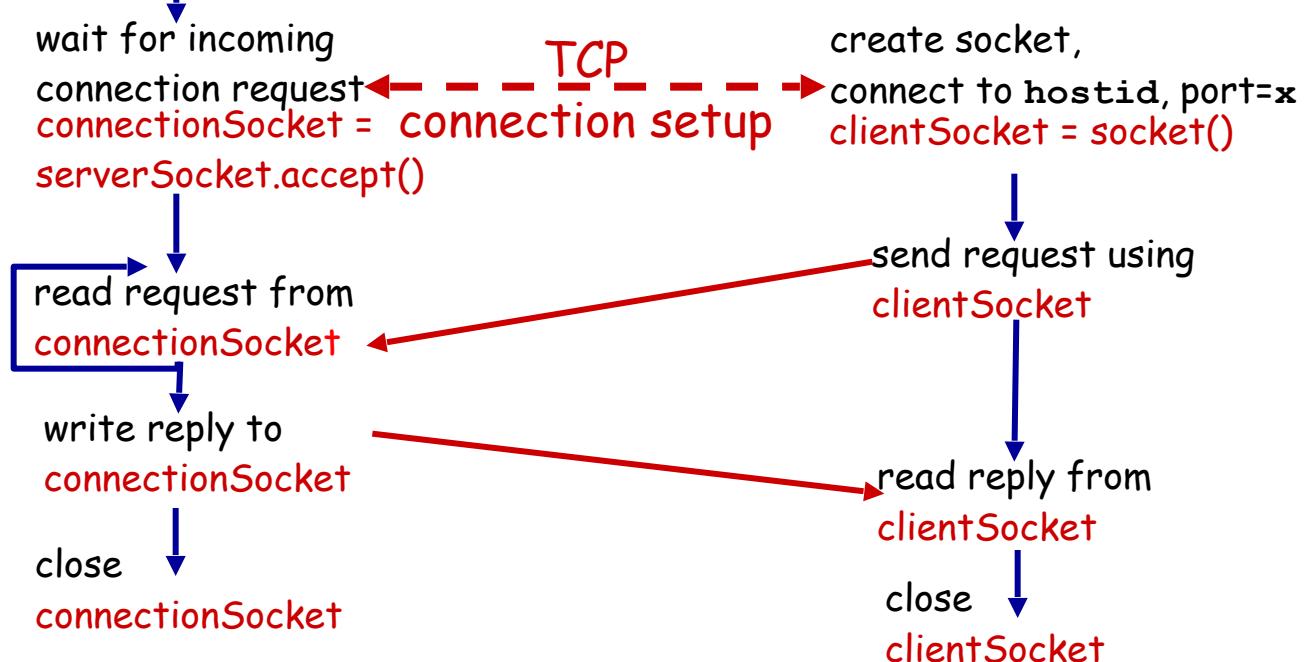
client

create socket,
connect to hostid, port=x
`clientSocket = socket()`

send request using
`clientSocket`

read reply from
`clientSocket`
close
`clientSocket`

TCP



Example app: TCP client

Python TCPClient

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```

create TCP socket for
server, remote port
12000

No need to attach server
name, port



Example app: TCP server

Python TCPServer

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET,SOCK_STREAM)
serverSocket.bind(("",serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
while True:
    connectionSocket, addr = serverSocket.accept()
    sentence = connectionSocket.recv(1024).decode()
    capitalizedSentence = sentence.upper()
    connectionSocket.send(capitalizedSentence.encode())
    connectionSocket.close()
```

create TCP welcoming
socket → serverSocket = socket(AF_INET,SOCK_STREAM)

server begins listening
for incoming TCP
requests → serverSocket.bind(("",serverPort))
loop forever → serverSocket.listen(1)

server waits on accept()
for incoming requests, new
socket created on return → print 'The server is ready to receive'

read bytes from socket
(but not address as in
UDP) → while True:

close connection to this
client (but not welcoming
socket) → connectionSocket, addr = serverSocket.accept()